

Data Structures and Algorithms

Instructor: Prof. Tianyi ZANG
School of Computer Science and Technology
Harbin Institute of Technology
tianyi.zang@gmail.com



		Insertion	Selection	Bubble	Shell	Merge	Heap	Quick	Quick3
Random									
Nearly Sorted									
Reversed									
Few Unique									

Sorting

Instructor: Prof. Tianyi ZANG
School of Computer Science and Technology
Harbin Institute of Technology
tianyi.zang@gmail.com



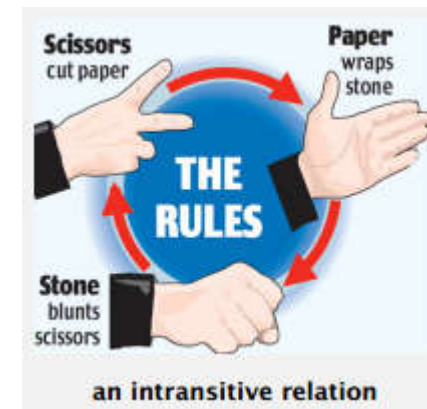
Sorting

- Sorting is one of the most fundamental algorithmic problems within computer science.
- It has been claimed that as many as 25% of all CPU cycles are spent sorting
- One reason why sorting is so important is that once a set of items is sorted, many other problems become easy.
- Sorting is fundamental to most other algorithmic problems, for example binary search.



What is sorting?

- It is the problem of taking an arbitrary permutation of n items and rearranging them into the total order.
- A **total order** is a binary relation \leq that satisfies properties below:
 - Antisymmetry: if $v \leq w$ and $w \leq v$, then $v = w$.
 - Transitivity: if $v \leq w$ and $w \leq x$, then $v \leq x$.
 - Totality: either $v \leq w$ or $w \leq v$ (comparable)
- Example:
 - Standard order for natural and real numbers.
 - Chronological order for dates or times.
 - Alphabetical order for strings
 - ...





Issues in Sorting

- Increasing or Decreasing Order?
 - The same algorithm can be used by both
 - all we need do is to change in the comparison function as we desire.
- What about equal keys?
 - Does the order matter or not?
 - Maybe we need to sort on secondary keys.
- We can assume a comparison function which depends on the application.
 - Compare (a, b) should return '<', '>', or '='.



How do you sort? Sort strategies

- There are several different ideas which lead to sorting algorithms:
- **Insertion** - putting an element in the appropriate place in a sorted list yields a larger sorted list.
- **Exchange** - rearrange pairs of elements which are out of order, until no such pairs remain.
- **Selection** - extract the largest element from the list, remove it, and repeat.
- **Distribution** - separate into piles based on the first letter, then sort each pile. (e.g. contact list in your mobile)
- **Merging** - Two sorted lists can be easily combined to form a sorted list.



Insertion Sort



Insertion Sort

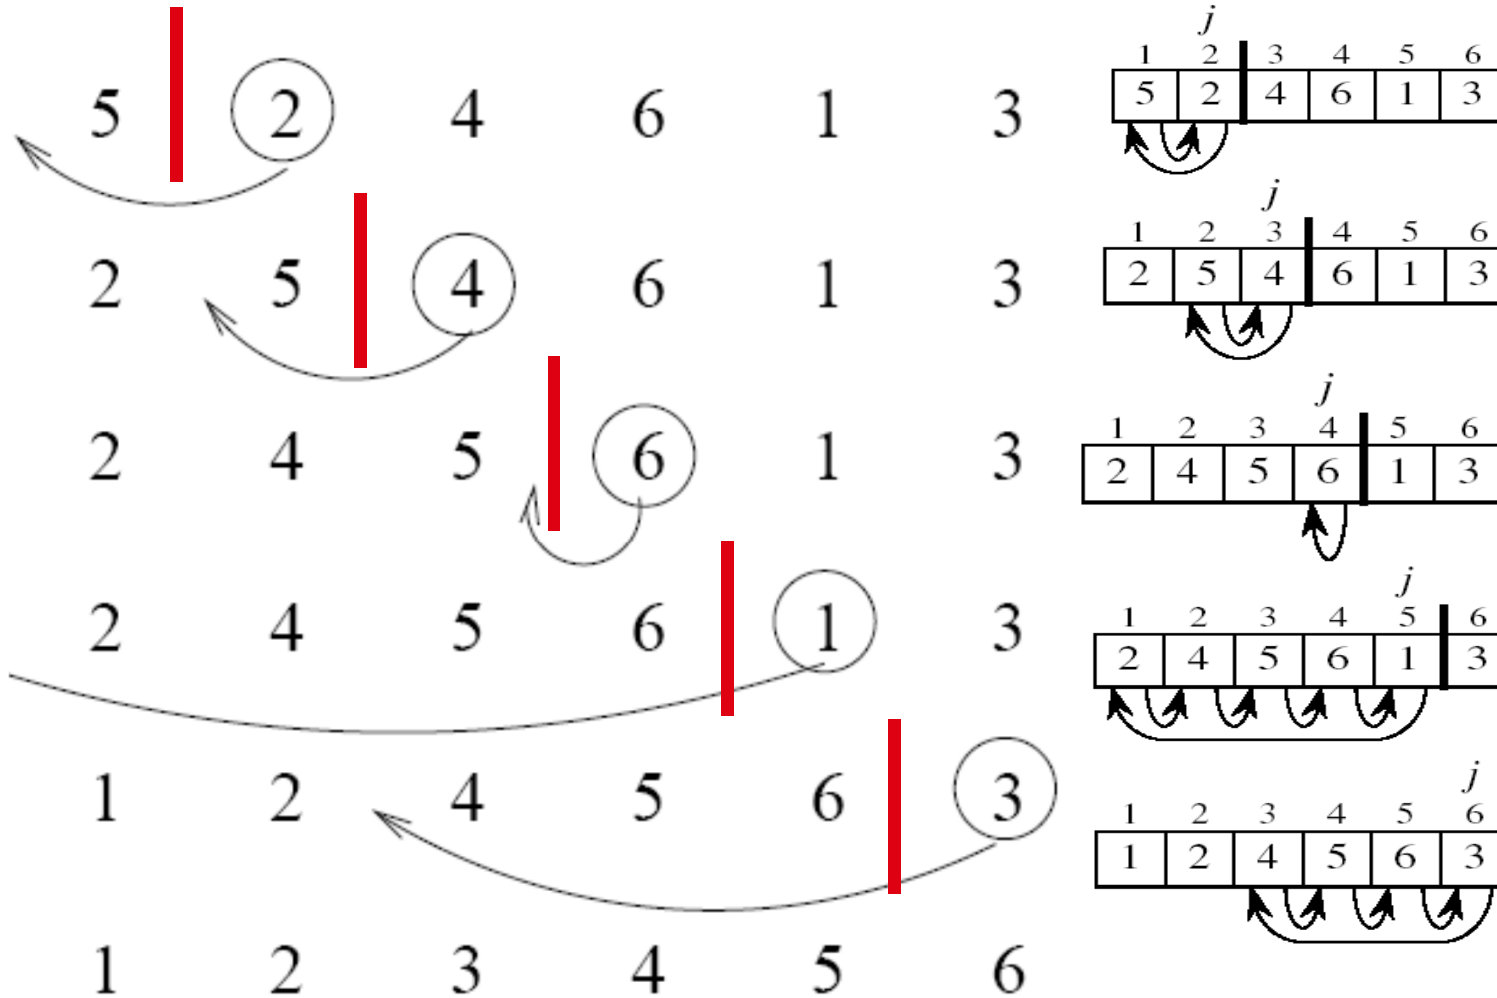
- In insertion sort, we repeatedly add elements to a sorted subset of our data, inserting the next element in order:

	i=1	2	3	4	5	6	7
42	20	17	13	13	13	13	13
20	42	20	17	17	14	14	14
17	17	42	20	20	17	17	15
13	13	13	42	28	20	20	17
28	28	28	28	42	28	23	20
14	14	14	14	14	42	28	23
23	23	23	23	23	23	42	28
15	15	15	15	15	15	15	42





Insertion Sort





Insertion Sort

- If the elements are in **an array** , we scan from bottom to top until we find the j such that

$$A[j] \leq x \leq A[j + 1]$$

then move from $j+1$ to the end down one to make room.

- If the elements are in **a linked list** , we do the sequential search until we find where the element goes,

then insert the element there.

No other elements need move!



Insertion Sort : Complexity

- The best, worst, and average cases for insertion sort all differ.
- We do not necessarily have to scan the entire sorted section of the array.
- **Best case :**
 - the element always gets inserted at the end, so we don't have to move anything, and only **compare against the last sorted element**.
 - We have **$(n-1)$ insertions**, each with exactly one comparison and no data moves per insertion!
- What is this best case permutation?
 - It is when the array or list is already sorted!



Worst Case Complexity

- **Worst case:**
 - The element always gets inserted at the front,
 - so all the sorted elements must be moved at each insertion.
 - The i^{th} insertion requires $(i-1)$ comparisons and moves

so:

$$\sum_{i=1}^n (i-1) = n(n-1)/2 = O(n^2)$$

- What is the worst case permutation?
 - When the array is sorted in reverse order.



Average Case Complexity

- Average Case :
 - If we were given a random permutation, the chances of the i^{th} insertion requiring $0, 1, 2, \dots, (i-1)$ comparisons are equal, and hence $1/i$.
 - The expected number of comparisons is for the i^{th} insertion is:
 - Summing up over all n keys,

$$\sum_{k=1}^i \frac{(k-1)}{i} = \frac{1}{i} \sum_{k=1}^{i-1} k = \frac{1}{i} \times \frac{i(i-1)}{2} = \frac{i-1}{2}$$

$$\sum_{i=1}^n \frac{i-1}{2} = \frac{1}{2} \left(\frac{n(n-1)}{2} + \frac{n}{2} \right) = \frac{n^2}{4} + O(n)$$



Insertion Sort: Implementation

```
template <typename E, typename Comp>
void inssort(E A[], int n) { // Insertion Sort
    for (int i=1; i<n; i++)    // Insert i'th record
        for (int j=i; (j>0) && (Comp::prior(A[j], A[j-1]))); j--)
            swap(A, j, j-1);
}
```

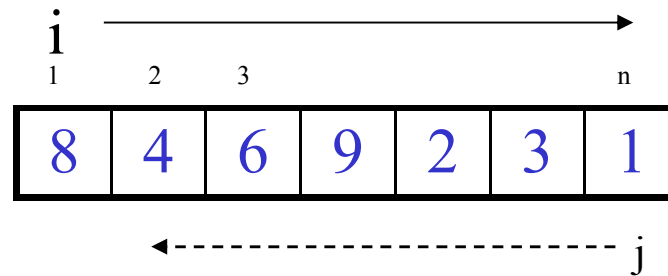


Bubble Sort



Bubble Sort

- Idea:
 - Repeatedly pass through the array
 - Swaps adjacent elements that are out of order



- Easier to implement, but slower than Insertion sort



Bubble Sort

- Simplest sorting algorithm
- **Steps:**
 1. Set flag = false
 2. Traverse the array and compare pairs of two elements
 - 2.1 If $E1 \leq E2$ - OK
 - 2.2 If $E1 > E2$ then
 - Swap($E1, E2$) and set flag = true
 3. If flag = true goto 1.



Bubble Sort: Example

8	4	6	9	2	3	1
---	---	---	---	---	---	---

$i = 1$ ←----- j

8	4	6	9	2	1	3
---	---	---	---	---	---	---

$i = 1$ ←----- j

8	4	6	9	1	2	3
---	---	---	---	---	---	---

$i = 1$ ←----- j

8	4	6	1	9	2	3
---	---	---	---	---	---	---

$i = 1$ ←----- j

8	4	1	6	9	2	3
---	---	---	---	---	---	---

$i = 1$ ←----- j

8	1	4	6	9	2	3
---	---	---	---	---	---	---

$i = 1$ j

1	8	4	6	9	2	3
---	---	---	---	---	---	---

$i = 1$ j

1	8	4	6	9	2	3
---	---	---	---	---	---	---

$i = 2$ j

1	2	8	4	6	9	3
---	---	---	---	---	---	---

$i = 3$ j

1	2	3	8	4	6	9
---	---	---	---	---	---	---

$i = 4$ j

1	2	3	4	8	6	9
---	---	---	---	---	---	---

$i = 5$ j

1	2	3	4	6	8	9
---	---	---	---	---	---	---

$i = 6$ j

1	2	3	4	6	8	9
---	---	---	---	---	---	---

$i = 7$

j



Bubble Sort: Implementation

```
template <typename E, typename Comp>
void bubsort(E A[], int n) { // Bubble Sort
    for (int i=0; i<n-1; i++)    // Bubble up i'th record
        for (int j=n-1; j>i; j--)
            if (Comp::prior(A[j], A[j-1]))
                swap(A, j, j-1);
}
```



Bubble Sort

- Demo (Bubble-sort with Hungarian (-Csángó-) folk dance) and
- consider its time complexity



Bubble Sort: Complexity

- One traversal = move the maximum element at the front
- Traversal #i: $n - i$ operations

- Running time:

$$(n - 1) + (n - 2) + \dots + 1 = (n - 1) n / 2 = O(n^2)$$

- Bubble Sort's running time is roughly the same in the best, average, and worst cases



Selection Sort



Selection Sort

- The most natural and easiest sorting algorithm is selection sort ,
 - where we repeatedly find the smallest element,
 - move it to the front, then repeat...

	i=0	1	2	3	4	5	6
42	13	13	13	13	13	13	13
20	<u>20</u>	14	14	14	14	14	14
17	17	<u>17</u>	15	15	15	15	15
13	42	42	<u>42</u>	17	17	17	17
28	28	28	28	<u>28</u>	20	20	20
14	14	20	20	20	<u>28</u>	23	23
23	23	23	23	23	23	<u>28</u>	28
15	15	15	17	42	42	42	<u>42</u>



Selection Sort : Implementation

```
template <typename E, typename Comp>
void selsort(E A[], int n) { // Selection Sort
    for (int i=0; i<n-1; i++) { // Select i'th record
        int lowindex = i;      // Remember its index
        for (int j=n-1; j>i; j--) // Find the least value
            if (Comp::prior(A[j], A[lowindex]))
                lowindex = j;    // Put it in place
        swap(A, i, lowindex);
    }
}
```

- If elements are **in an array**,
 - swap the first with the smallest element- thus only one array is necessary.
- If elements are **in a linked list**,
 - we must keep two lists, one sorted and one unsorted, and
 - always add the new element to the back of the sorted list.



Selection Sort : Complexity

- The best case, worst case, and average cases are all the same!
- To find the largest takes (n-1) steps, to find the second largest takes (n-2) steps, ..., so, takes totally

$$\sum_{i=1}^{n-1} i = n(n-1)/2 = n^2/2 + n/2 = O(n^2)$$



Summary: exchange sorts

	Insertion	Bubble	Selection
Comparisons:			
Best Case	$\Theta(n)$	$\Theta(n^2)$	$\Theta(n^2)$
Average Case	$\Theta(n^2)$	$\Theta(n^2)$	$\Theta(n^2)$
Worst Case	$\Theta(n^2)$	$\Theta(n^2)$	$\Theta(n^2)$
Swaps:			
Best Case	0	0	$\Theta(n)$
Average Case	$\Theta(n^2)$	$\Theta(n^2)$	$\Theta(n)$
Worst Case	$\Theta(n^2)$	$\Theta(n^2)$	$\Theta(n)$

- Swapping adjacent records is called an exchange.
- Thus, these sorts are sometimes referred to as **exchange sorts**.



Why exchange sorts are slow?

- The cost of any exchange sort can be at best the **total number of steps** that the records in the array must **move to reach their “correct” location** (i.e., the number of inversions for each record).
- **Inversion:** a pair (i, j) such that $i < j$, but $\text{Array}[i] > \text{Array}[j]$
- Array of size N
 - average number of inversions in a random set of elements is $N(N-1)/4$
- Exchange Sort only swaps adjacent elements
 - only removes 1 inversion!



Shellsort



Shellsort (diminishing increment sort)

- Shellsort is also known as *diminishing increment sort*.
- Invented by Donald Shell in 1959.
- 1st algorithm to **break the quadratic time** barrier but few years later, a sub quadratic time bound was proven
- Shellsort works by comparing elements that are **distant rather than adjacent elements** in an array.



Shellsort (diminishing increment sort)

- The **distance between comparisons decreases** as the sorting algorithm runs until the last phase in which adjacent elements are compared



Shellsort (diminishing increment sort)

- Shellsort makes comparisons and swaps between **non-adjacent** elements
- Shellsort 's strategy is to make the list “mostly sorted ” so that a final Insertion Sort can finish the job
- Idea
 - Shellsort breaks the array of elements into “virtual ” sublists according to the increment.
 - Each sublist is sorted using an Insertion Sort.
 - Another group of sublists is then chosen and sorted,
 - and so on.



Example demo

	1	2	3	4	5	6	7	8	9
Initial list	40	25	49	25*	16	21	08	30	13
d = 4	40	25	49	25*	16	21	08	30	13
	13	21	08	25*	16	25	49	30	40
d = 2	13	21	08	25*	16	25	49	30	40
	08	21	13	25*	16	25	40	30	49
d = 1	08	21	13	25*	16	25	40	30	49
	08	13	16	21	25*	25	30	40	49



Shellsort: Implementation

```
// Modified version of Insertion Sort for varying increments
template <typename E, typename Comp>
void inssort2(E A[], int n, int incr) {
    for (int i=incr; i<n; i+=incr)
        for (int j=i; (j>=incr) &&
                (Comp::prior(A[j], A[j-incr])); j-=incr)
            swap(A, j, j-incr);
}

template <typename E, typename Comp>
void shellsort(E A[], int n) { // Shellsort
    for (int i=n/2; i>2; i/=2) // For each increment
        for (int j=0; j<i; j++) // Sort each sublist
            inssort2<E, Comp>(&A[j], n-j, i);
    inssort2<E, Comp>(A, n, 1);
}
```



Shellsort: Complexity

- The average-case performance of Shellsort (for “divisions by three ” increments) is $O(n^{1.5})$.
- Shellsort is substantially better than Insertion Sort, or any of exchange sorts ($O(n^2)$).

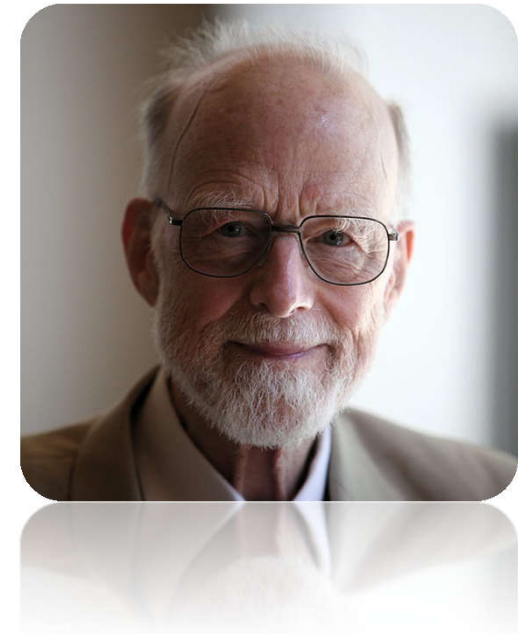


Quicksort



Quicksort

- Sir Charles Antony Richard Hoare (born 11 January 1934), commonly known as Tony Hoare or C. A. R. Hoare
- He developed the sorting algorithm Quicksort in 1960 at age 26.
- He also developed Hoare logic for verifying program correctness, and the formal language Communicating Sequential Processes (CSP)
- Won ACM Turing Award for "fundamental contributions to the definition and design of programming languages" on 27 October 1980
- For many years under his leadership his Oxford department worked on formal specification languages such as CSP and Z.





Quicksort

- Quicksort when properly implemented, is the fastest known general-purpose in-memory sorting algorithm in the average case.
- It does not require the extra array needed by Mergesort, so it is space efficient as well.



Quicksort t-shirt



```
public static void quicksort(char[] items, int left, int right)
{
    int i, j;
    char x, y;

    i = left; j = right;
    x = items[(left + right) / 2];

    do
    {
        while ((items[i] < x) && (i < right)) i++;
        while ((x < items[j]) && (j > left)) j--;

        if (i <= j)
        {
            y = items[i];
            items[i] = items[j];
            items[j] = y;
            i++; j--;
        }
    } while (i <= j);

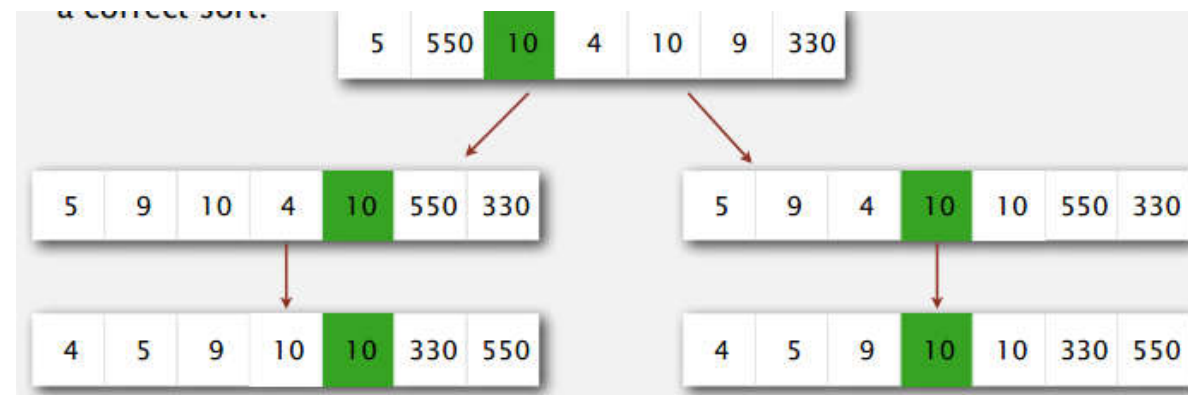
    if (left < j) quicksort(items, left, j);
    if (i < right) quicksort(items, i, right);
}
```



Partitioning - the heart of Quicksort

Partitioning

- To partition an array $a[]$ on element $x = a[i]$ is to rearrange it such that
 - x moves to position j (may be the same as i)
 - All entries to the left of x are $\leq x$.
 - All entries to the right of x are $\geq x$.
- Observations.
 - O1. After partitioning, x is 'in place.' No need to move x in order to reach a correct sort.





Partitioning - the heart of Quicksort

Partitioning

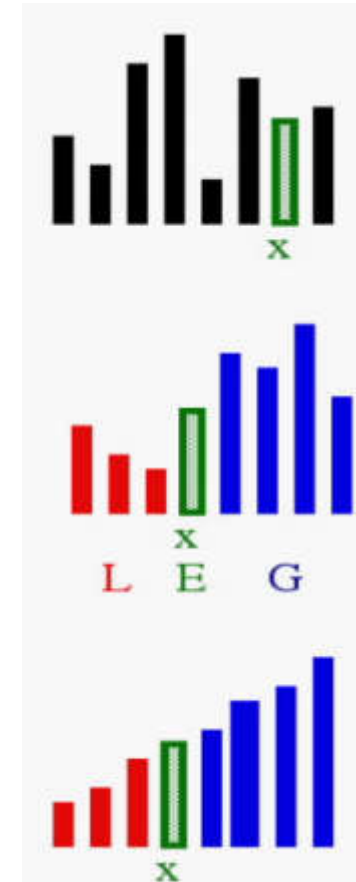
- Observations.
 - O1. After Partitioning , x is ‘in place.’ No need to move x in order to reach a correct sort.
 - O2. To Partition on a new pivot y, there’s no need to look beyond the confines of previously partitioned items.





Quicksort: Idea

- 1) Select: pick an element
- 2) Divide:
rearrange elements so that x goes to
its final position E
- 3) Recurse and Conquer:
recursively sort





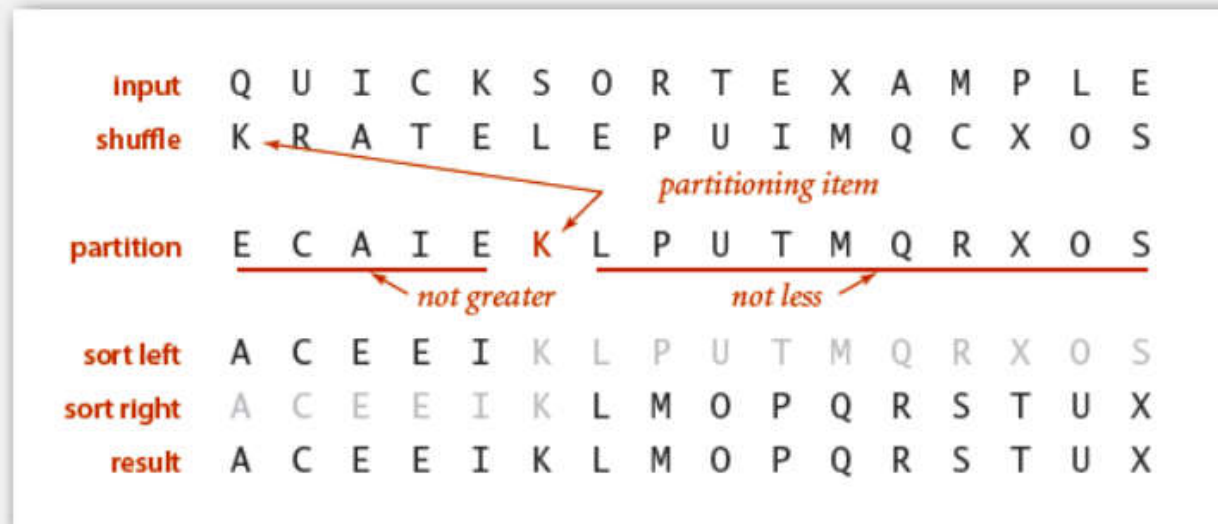
Quicksort: Idea

Basic plan.

- **Shuffle** the array.
- **Partition** so that, for some j
 - entry $a[j]$ is in place
 - no larger entry to the left of j
 - no smaller entry to the right of j
- **Sort** each piece recursively.



Sir Charles Antony Richard Hoare
1980 Turing Award





Quicksort partitioning



Basic plan.

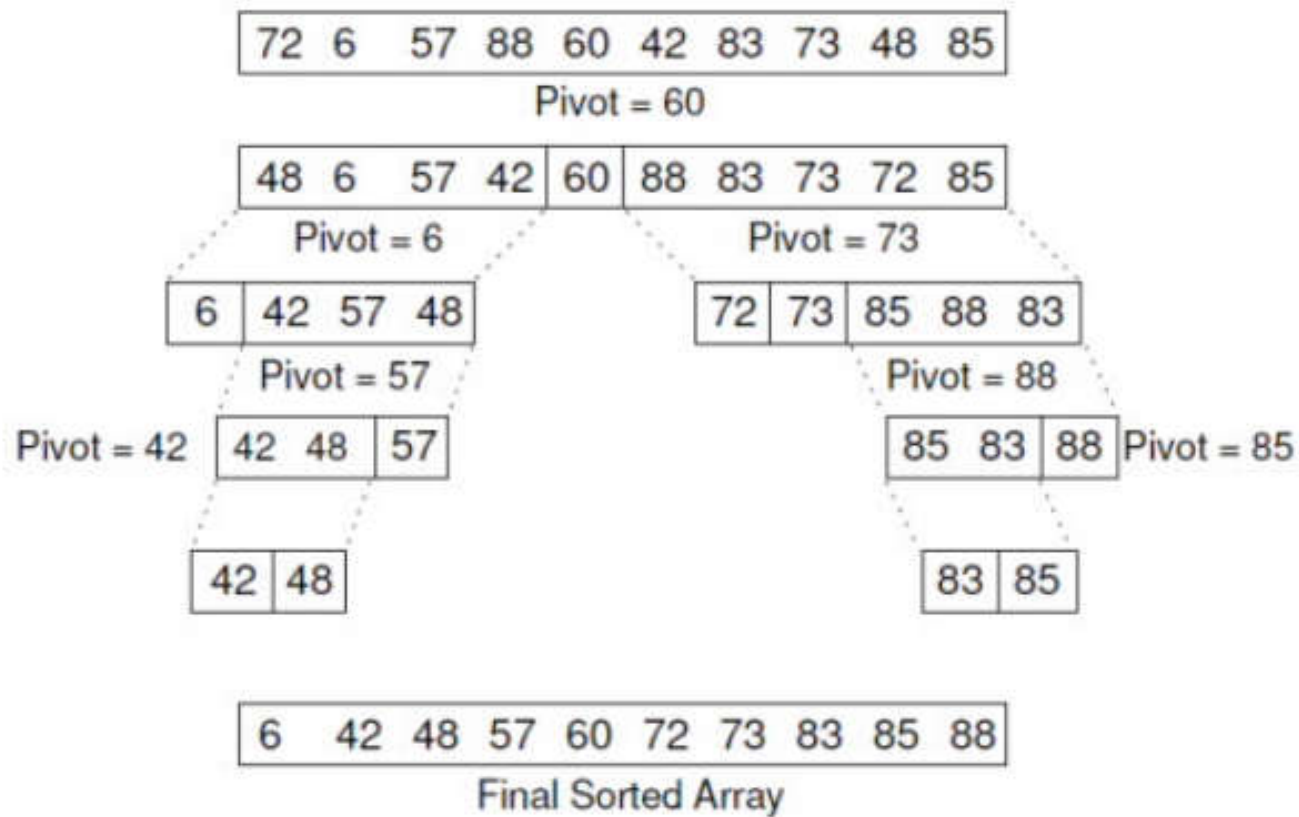
- Scan i from left for an item that belongs on the right.
- Scan j from right for an item that belongs on the left.
- Exchange $a[i]$ and $a[j]$.
- Repeat until pointers cross.

	i	j	a[i]															
			0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Initial values	0	16	K	R	A	T	E	L	E	P	U	I	M	Q	C	X	O	S
scan left, scan right	1	12	K	R	A	T	E	L	E	P	U	I	M	Q	C	X	O	S
exchange	1	12	K	C	A	T	E	L	E	P	U	I	M	Q	R	X	O	S
scan left, scan right	3	9	K	C	A	T	E	L	E	P	U	I	M	Q	R	X	O	S
exchange	3	9	K	C	A	I	E	L	E	P	U	T	M	Q	R	X	O	S
scan left, scan right	5	6	K	C	A	I	E	L	E	P	U	T	M	Q	R	X	O	S
exchange	5	6	K	C	A	I	E	E	L	P	U	T	M	Q	R	X	O	S
scan left, scan right	6	5	K	C	A	I	E	E	L	P	U	T	M	Q	R	X	O	S
final exchange	6	5	E	C	A	I	E	K	L	P	U	T	M	Q	R	X	O	S
result	6	5	E	C	A	I	E	K	L	P	U	T	M	Q	R	X	O	S

Partitioning trace (array contents before and after each exchange)



Quicksort: Example





Quicksort Implementation

```
int FindPivot( int i, int j ) /* A is an extra array */
/* if A[i],...A[j] are the same, then return 0;
else, the bigger one. */
{  keytype  firstkey = A[i].key ; /* */
   int k ;
   for ( k=i+1 ; k<=j; k++ ) /*
       if ( A[k].key > firstkey )
           return k ;
       else if ( A[k].key < firstkey )
           return i ;
   return 0 ;
}
```



Quicksort Implementation

```
int Partition ( int i , int j , keytype pivot )  
/* partition A[i],...,A[j]*/  
{  int l , r ;  
    do{  
        for( l = i ; A[l].key < pivot ; l++ ) ;  
        for( r = j ; A[r].key >= pivot ; r-- ) ;  
        if( l < r ) swap(A[l],A[r]);  
    } while( l <= r );  
    return  l ;  
}
```



Quicksort Implementation

```
void QuickSort ( int i , int j )
{   keytype pivot;
    int k; //
    int pivotindex ; //
    pivotindex = FindPivot ( i , j );
    if( pivotindex != 0 ) { //
        pivot=A[pivotindex].key;
        k=Partition ( i , j , pivot );
        QuickSort ( i , k-1);
        QuickSort ( k , j );
    }
}
```



Partitioning - the heart of Quicksort

- Demo 2: comparison bubble sort and quick sort (Visualization of Quick sort.mp4)
- And think about the time complexity of quicksort.



Time complexity

- Best case
 - Each partition, the length of the sub-lists are the same.

$$T(n) \leq 2T(n/2) + n$$

$$\leq 2(2T(n/4) + n/2) + n = 4T(n/4) + 2n$$

$$\leq 4(2T(n/8) + n/4) + 2n = 8T(n/8) + 3n$$

... ..

$$\leq nT(1) + n\log_2 n = O(n\log_2 n)$$

- *Time complicity* $O(n^2)$
- *Time complicity* $O(n)$



Time complexity

- Best case
 - Each partition, the length of the sub-lists are the same.

$$T(n) \leq 2T(n/2) + n$$

$$\leq 2(2T(n/4) + n/2) + n = 4T(n/4) + 2n$$

$$\leq 4(2T(n/8) + n/4) + 2n = 8T(n/8) + 3n$$

... ..

$$\leq nT(1) + n\log_2 n = O(n\log_2 n)$$

- *Time complicity* $O(n\log_2 n)$
- *Space complicity* $O(n)$



Time complexity

- Worst case
 - when the pivot does not divide the sequence in two
 - At each step, the length of the sequence is only reduced by 1
 - Total running time

$$\sum_{i=1}^{n-1} (n-i) = \frac{1}{2} n(n-1) = O(n^2)$$

- Time complicity $O(n^2)$
- Space complicity $O(n)$



Time complexity

- General/Average case:
 - Time spent at level i in the tree is $O(n)$
 - Running time: $O(n) * O(\text{height})$
 - Time complexity $O(n \log_2 n)$
 - Spatial complexity $O(n)$



Mergesort



Mergesort

- Mergesort is one of the simplest sorting algorithms conceptually, and
- has good performance both in the asymptotic sense and in empirical running time.
- Basic idea is to merge two or more ordered lists into one ordered list.



Merge order



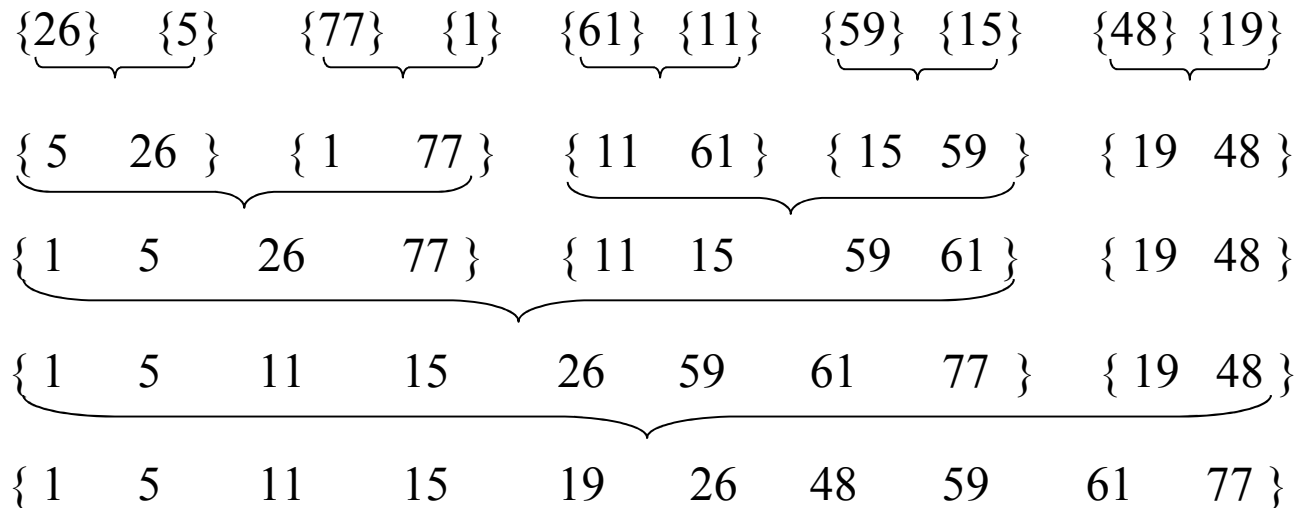
```
void Merge (int s , int m , int t , LIST A , LIST B)
/*Merge ordered lists of A[s],...,A[m] and A[m+1],...,A[t] into ordered one of
  B[s],...,B[t]*/
{   int i = s ; j = m+1 , k = s ;//initialization
    /* if not empty, place smaller into B[k] */
    while ( i <= m && j <= t )
        B[k++] = ( A[ i ].key <= A[ j ].key) ? A[i++] : A[j++] ;
    /* if the first list is not empty, copy the left into B */
    while ( i <= m )   B[k++] = A[i++] ;
    /* if the second list is not empty, copy the left into B */
    while ( j <= t )   B[k++] = A[j++] ;
}/*time complex. :  $O(t - s + 1)$  ; space complex. :  $O(t - s + 1)$  */
```





Merge order: non recursive

- **Idea of two-way mergesort**
 - Firstly regard the list of n records as n lists of 1 record
 - Merging each two lists, we have $\lceil n/2 \rceil$ ordered lists of 2 records
 - Merging each two lists, we have $\lceil n/4 \rceil$ ordered lists of 4 records
 -
 - At last, we have a ordered list of n records.





Merge order: non recursive algorithm

Idea for one pass: merge the two adjacent lists which length is h in new ordered list which length is $2h$

```
void MergePass (int n , int  $h$  , LIST A , LIST B)
```

```
{   int i , t ;
```

```
    for ( i=1 ;  $i+2*h-1 \leq n$  ;  $i+=2*h$  )
```

```
        Merge(i,  $i+h-1$ ,  $i+2*h-1$ , A, B) ; /*merge two ordered lists of length  $h$  */
```

```
    if (  $i+h-1 < n$  ) /* there are still two lists left, but the length of the second one is less than  $h$  */
```

```
        Merge( i,  $i+h-1$ ,  $n$  , A, B) ; /* merge the last two lists */
```

```
    else /* if  $i \leq n$  and  $i+h-1 > n$ , only one list left, make a copy */
```

```
        for ( t= i ; t<= n ; t++ )
```

```
            B[t] = A[t] ;
```

```
} /* Mpass */
```



Merge order: non recursive algorithm

Idea for Control the loop of merge: the length of merged list reach to n

```
void MergeSort ( int n , LIST A )  
{ /* two-way merge sort */  
  int  $h = 1$  ; /* the length of merging list, initialized as 1 */  
  LIST B ;  
  while ( $h < n$ ){  
    MergePass (n ,  $h$  , A , B) ;  
     $h = 2 * h$  ;  
    MergePass (n ,  $h$  , B , A) ; /* swap A and B */  
     $h = 2 * h$  ;  
  }  
} /* MergeSort */
```



Performance analysis

- Time complexity:
 - Each pass $O(n)$ times merge are needed
 - $\lceil \log_2 n \rceil$ passed merge are needed for all.
 - $O(n \log_2 n)$ for the best, worst and average case.
- Space complexity:
 - $O(n)$



- merge sort vs quick sorts



Merge sort: recursion

Idea: **divide-and-conquer**

- **Divide**: If S has at least two elements (nothing needs to be done if S has zero or one elements),
 - remove all the elements from S and
 - put them into two sequences, S_1 and S_2 , each containing about half of the elements of S . i.e.
 - S_1 contains the first $n/2$ elements and
 - S_2 contains the remaining $n/2$ elements.
- **Recur** : Recursive sort sequences S_1 and S_2
- **Conquer**: Put back the elements into S by merging the sorted sequences S_1 and S_2 into a unique sorted sequence



Merge sort: recursive algorithm

```
void MergeSort ( LIST A , LIST B , int low , int high )  
/* divide-and-conquer, mergesort A[low], ..., A[high] */  
{   int mid = (low+high)/2 ;  
    if (low<high){ /* high-low>0 */  
        MergeSort (A , B , low , mid) ;  
        MergeSort (A , B , mid+1 , high) ;  
        Merge (low , mid , high , A , B) ;  
    }  
} /* MergeSort */
```

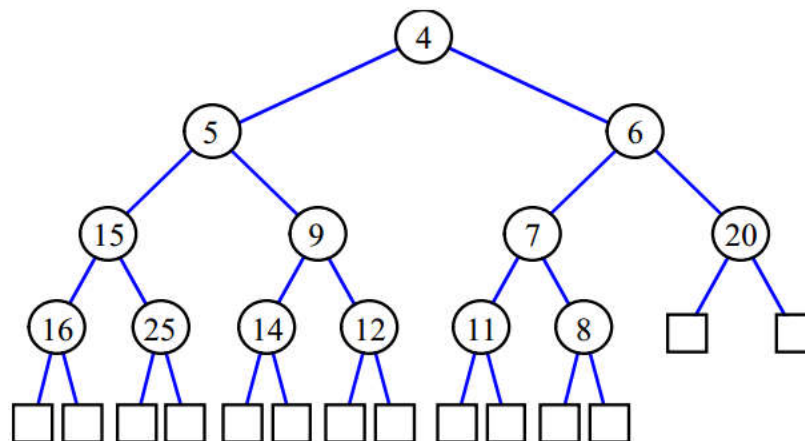


Heap sort



Heaps

- A heap is a binary tree T that stores a key-element pairs at its internal nodes
- It satisfies two properties:
 - MinHeap: $\text{key}(\text{parent}) \leq \text{key}(\text{child})$
 - [OR MaxHeap: $\text{key}(\text{parent}) \geq \text{key}(\text{child})$]
 - all levels are full, except the last one, which is left-filled
- A heap T storing n keys has height $h = \log(n + 1)$





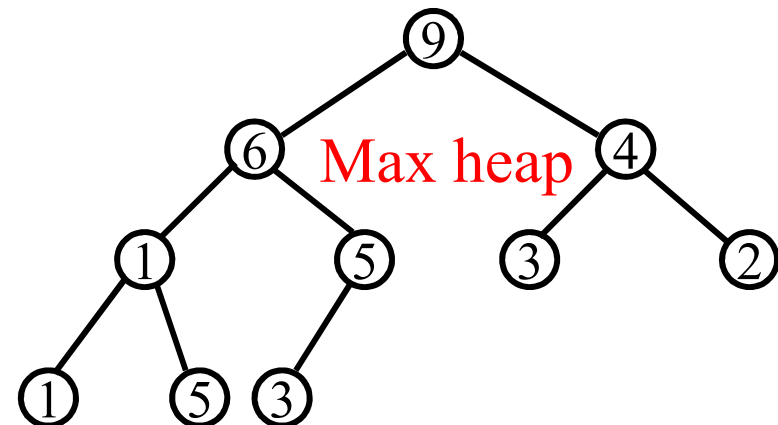
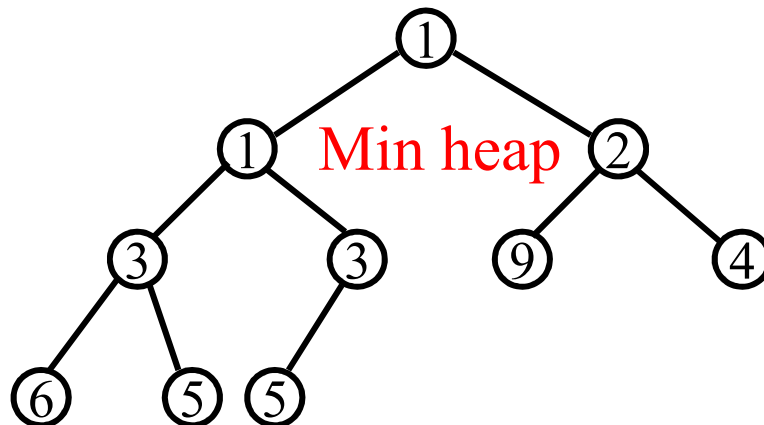
Properties of Heaps

MinHeap: representation of array A

- (1) if $2*i \leq n$, then $A[i].key \leq A[2*i].key$
- (2) if $2*i+1 \leq n$, then $A[i].key \leq A[2*i+1].key$

MaxHeap: representation of array A

- (1) if $2*i \leq n$, then $A[i].key \geq A[2*i].key$
- (2) if $2*i+1 \leq n$, then $A[i].key \geq A[2*i+1].key$

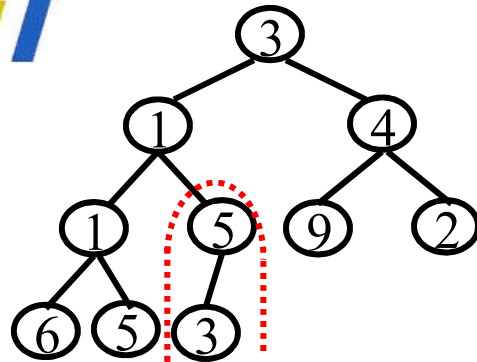




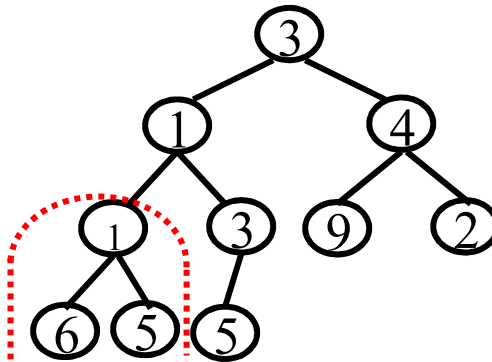
Heap sort

Steps :

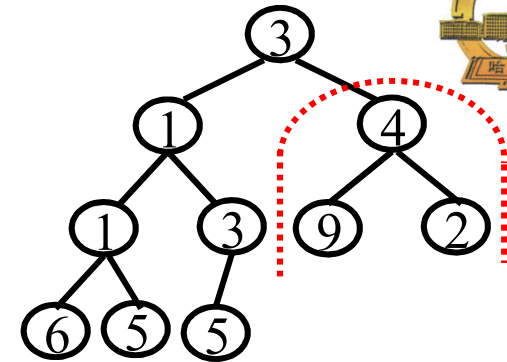
- **Initialize** the array of complete binary tree with unsorted records
- **Construct heaps**: let $i = n/2, \dots, 2, 1$, construct heap with root 1 by $\text{PushDown}(i, n)$.
- **Heap sort**: let $i = n, n-1, \dots, 2$
 1. **swap**: swap the root (the least value) with i element (the largest value), i.e. $\text{swap}(A[1], A[i])$;
 2. **manage**: make the left $i-1$ elements a heap by $\text{PushDown}(1, i-1)$;
 3. **repeat**: until the sequence of $A[1], A[2], \dots, A[n]$ is ordered.



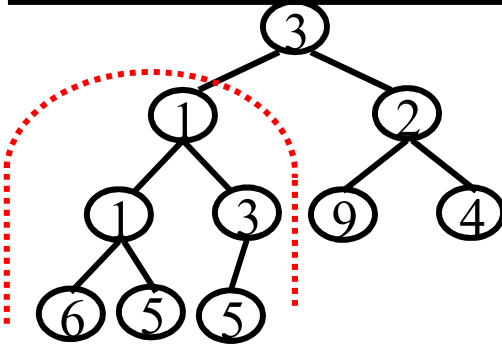
3	1	4	1	5	9	2	6	5	3
---	---	---	---	---	---	---	---	---	---



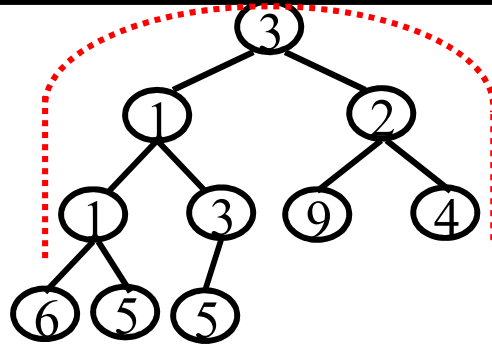
3	1	4	1	3	9	2	6	5	5
---	---	---	---	---	---	---	---	---	---



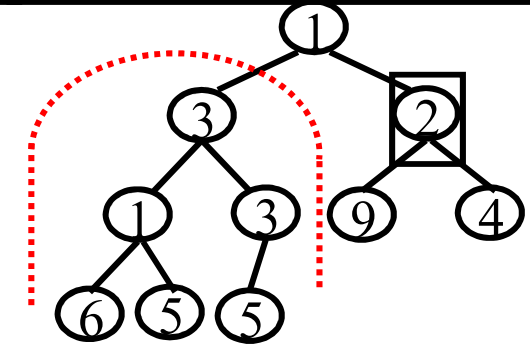
3	1	4	1	3	9	2	6	5	5
---	---	---	---	---	---	---	---	---	---



3	1	2	1	3	9	4	6	5	5
---	---	---	---	---	---	---	---	---	---

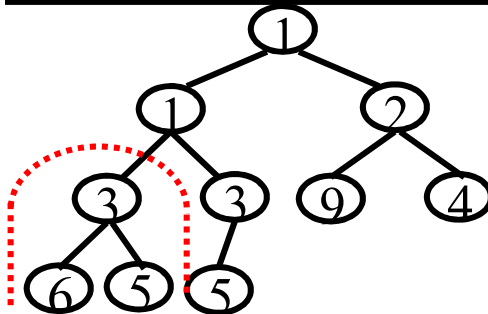


3	1	2	1	3	9	4	6	5	5
---	---	---	---	---	---	---	---	---	---

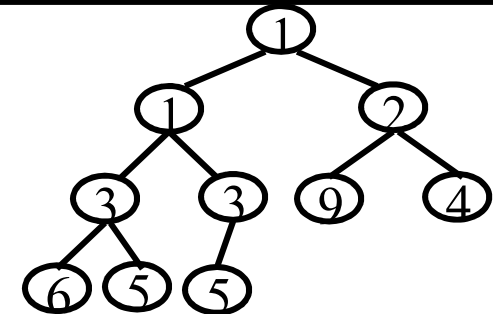


1	3	2	1	3	9	4	6	5	5
---	---	---	---	---	---	---	---	---	---

The process of
heap construction



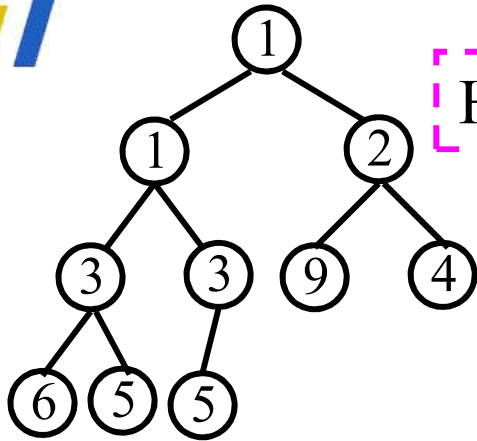
1	1	2	3	3	9	4	6	5	5
---	---	---	---	---	---	---	---	---	---



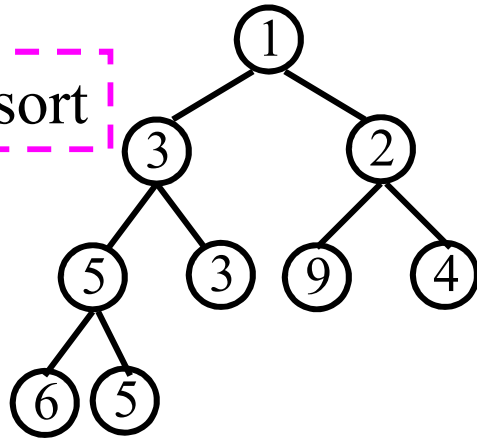
1	1	2	3	3	9	4	6	5	5
---	---	---	---	---	---	---	---	---	---



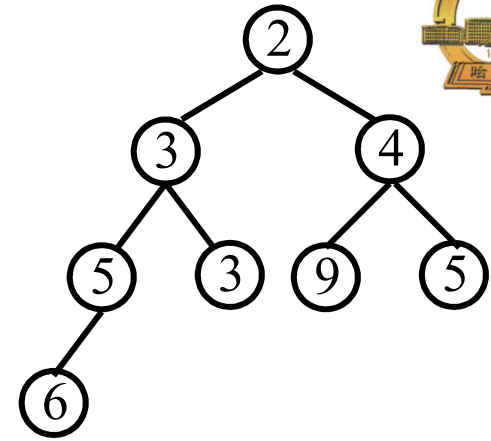
Heap sort



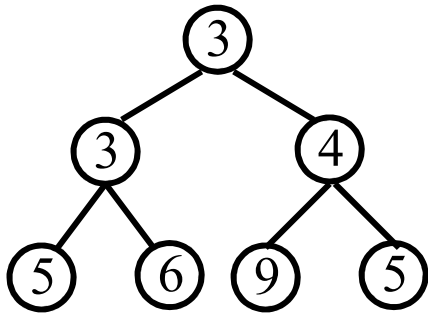
1	1	2	3	3	9	4	6	5	5
---	---	---	---	---	---	---	---	---	---



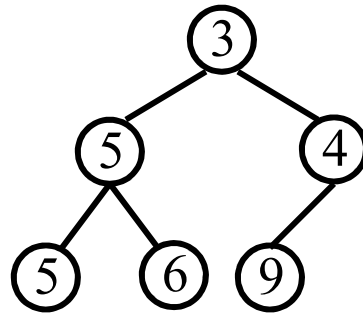
1	3	2	5	3	9	4	6	5	1
---	---	---	---	---	---	---	---	---	---



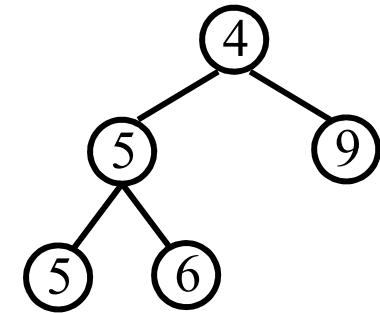
2	3	4	5	3	9	5	6	1	1
---	---	---	---	---	---	---	---	---	---



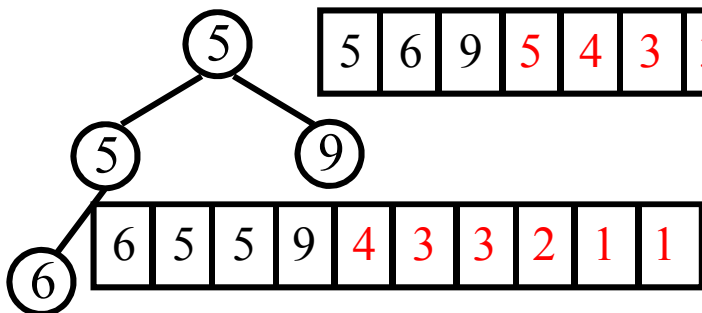
3	3	4	5	6	9	5	2	1	1
---	---	---	---	---	---	---	---	---	---



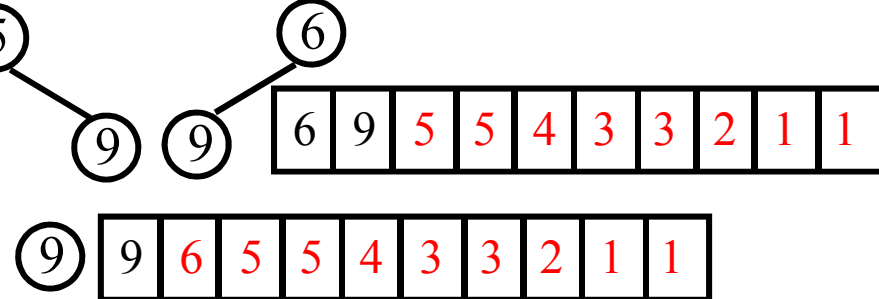
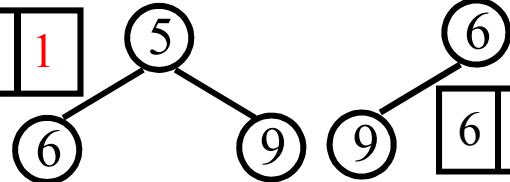
3	5	4	5	6	9	3	2	1	1
---	---	---	---	---	---	---	---	---	---



4	5	9	5	6	3	3	2	1	1
---	---	---	---	---	---	---	---	---	---



5	6	9	5	4	3	3	2	1	1
---	---	---	---	---	---	---	---	---	---



6	9	5	5	4	3	3	2	1	1
---	---	---	---	---	---	---	---	---	---

9	9	6	5	5	4	3	3	2	1
---	---	---	---	---	---	---	---	---	---



Implementation of heap sort

```
void HeapSort ( int n , LIST A )
{   int i;
    for( i=n/2; i>=1; i--) /*construct heap from the most right non-leaf node*/
        PushDown( i, n);
        /*make it a heap, i is the root, n is the index of max leaf node */
    for( i=n; i>=2; i--) {
        swap(A[1],A[i]); // swap root with i which is index of max leaf node
        PushDown( 1, i-1 );
        /*make a heap, 1 is root*/
    }
}
```



- 20151225 Friday week 15
- The course on Sunday is changed to 3-4 on Monday, 28 Dec.
- The last course for the presentation is arranged 3-4 on Thursday, 31 Dec.



Implementation of heap sort



```
void PushDown(int first,int last)
{
    /*manage heap: push A[first] down to a proper position*/
    int r=first; /* r is the right position, initialized as first*/
    while(r<=last/2) /* A[r] is not leaf, otherwise it is a heap */
        if((r==last/2) && (last%2==0)) { /* r has one left son of 2*r*/
            if(A[r].key>A[2*r].key)
                swap(A[r],A[2*r]);/*push down, swap*/
            r=last; /*A[r].key <=A[2*r].key or >=, if reach the leaf, the loop stop*/
        } else if((A[r].key>A[2*r].key)&&(A[2*r].key<=A[2*r+1].key)) {
            /*root > left son and left son <=right son (select smaller one)*/
            swap(A[r],A[2*r]); /*swap with the left son*/
            r=2*r; /* push down to the position*/
        } else if((A[r].key>A[2*r+1].key)&&(A[2*r+1].key<A[2*r].key)) {
            /* root > right son, and the right son < left son (select smaller one) */
            swap(A[r],A[2*r+1]); /*swap with the right son*/
            r=2*r+1; /*push down to the position*/
        } else /*A[r] is the heap */
            r=last;
    } /*PushDown*/
}
```



Performance of heap sort

Performance analysis

- Time complexity of PushDown:
 - Within a while, $r = \text{first} * 2^i > \text{last}/2$, i.e. $i > \log_2(\text{last}/\text{first}) - 1$
 - $O(\log(\text{last}/\text{first})) = O(\log_2 n)$
- Time complexity of HeapSort :
 - $O(n \log_2 n)$ for best, worst and average cases.



How do you sort? Sort strategies

- There are several different ideas which lead to sorting algorithms:
- Insertion - putting an element in the appropriate place in a sorted list yields a larger sorted list.
- Exchange - rearrange pairs of elements which are out of order, until no such pairs remain.
- Selection - extract the largest element from the list, remove it, and repeat.
- **Distribution** - separate into piles based on the first letter, then sort each pile. (e.g. contact list in your mobile)
- Merging - Two sorted lists can be easily combined to form a sorted list.



Radix sort



Review of the performance of sorting algorithms

- **Review:** summary of the performance of sorting algorithms
Frequency of execution of instructions in the inner loop:

algorithm	average	extra space	operations on keys
insertion sort	n^2	no	compareTo()
selection sort	n^2	no	compareTo()
mergesort	$n \lg n$	n	compareTo()
quicksort	$1.39 n \lg n$	$c \lg n$	compareTo()

- **lower bound:** $n \lg n - 1.44n$ compares are required by any algorithm
//Based on the Stirling's approximation: $\log_2 n! \approx n \log_2 n - 1.44n + O(\log_2 n)$
- **Q:** Can we do better (despite the lower bound)?



Radix sort

- We can do better (despite the lower bound)
 - Yes, if we do not depend on comparisons
 - Linear time complexity $O(n)$
- Idea of radix sort (bin sort, bucket sort)
 - Radix sort sorts integers by processing individual digits, by comparing individual digits sharing the same significant position.
 - Radix sort is not limited to integers.
 - Because integers can represent strings of characters (e.g., names or dates) and specially formatted floating point numbers.
- Least-significant-digit (LSD) radix sort
- Most-significant-digit (MSD) radix sorts



Radix sort

Ideas:

- There is a one-to-one correspondence between the number of buckets and the number of values that can be represented by a digit.
- Assuming there is sequence of integers, which have *figure* of digits,
- Each digit ranges 0,1,..., to 9, i.e. the radix is 10

1. **Initializaton**: there are ten empty queues, $Q[0], Q[1], \dots, Q[9]$.

2. **Assignment**: for the *pass* pass, take all data in A, and put them in corresponding bucket according to *the most right pass digit (r)* of data.key, insert data into $Q[r]$.

/*The processing of the keys *begins at the least significant digit* (i.e., the *rightmost* digit), and *proceeds to the most significant digit* (i.e., the *leftmost* digit) */

3. **Gathering**: starting from $Q[0]$, take the all numbers from $Q[0], Q[1], \dots, Q[9]$, and insert them into array A in the taken order.

4. **Repeat step** 1,2,3, till *figure* passes, we have the array A ordered finally.



Demonstration of radix sort

321 986 123 432 543 018 765 678 987 789 098 890 109 901 210 012

Q[0]:890 210

Q[1]:321 901

Q[2]:432 012

Q[3]:123 543

Q[4]:

Q[5]:765

Q[6]:986

Q[7]:987

Q[8]:018 678 098

Q[9]:789 109

Q[0]:901 109

Q[1]:210 012 018

Q[2]:321 123

Q[3]:432

Q[4]:543

Q[5]:

Q[6]:765

Q[7]:678

Q[8]:986 987 789

Q[9]:890 098

Q[0]:012 018 098

Q[1]:109 123

Q[2]:210

Q[3]:321

Q[4]:432

Q[5]:543

Q[6]:678

Q[7]:765 789

Q[8]:890

Q[9]: 901 986 987

890 210 321 901 432 012 123 543 765 986 987 018 678 098 789 109

901 109 210 012 018 321 123 432 543 765 678 986 987 789 890 098

012 018 098 109 123 310 321 432 543 678 765 789 890 901 986 987



Implementation of the algorithm

```
void RadixSort( int figure, QUEUE &A)
{   QUEUE Q[10]; records data ;
    int pass, r, i ;
    for ( pass=1; pass<=figure ; pass++ ){
        for ( i=0 ; i<=9 ; i++ ) /*empty queue*/
            MAKENULL( Q[i] ) ;
        while ( !EMPTY( A ) ) /* assignment*/
            data = FRONT ( A ) ;
            DEQUEUE ( A ) ;
            r = Radix(data.key, pass) ;
            ENQUEUE( data , Q[r] ) ; }
    for ( i=0 ; i <=9 ; i++ ) /*gather*/
        while ( !EMPTY( Q[i] ) ) {
            data = FRONT ( Q[i] ) ;
            DEQUEUE( Q[i] ) ;
            ENQUEUE( data, A );}
}
```

```
/*return the pth digit of integer k */
int Radix ( int k, int p)
{ int power= 1 ;
  for ( int i=1; i<=p-1 ; i++ )
      power = power * 10 ;
  return
  (( k%(power*10))/power) ;
}
```

```
for (i=0;i<=9;i++) {
    Concatenate(Q[1], Q[i]);
    A=Q[0];
}/*connect queues together */
```



Implementation of the algorithm

If not sure the length of $Q[0]$, $Q[1]$,..., $Q[9]$, then we use linked queue

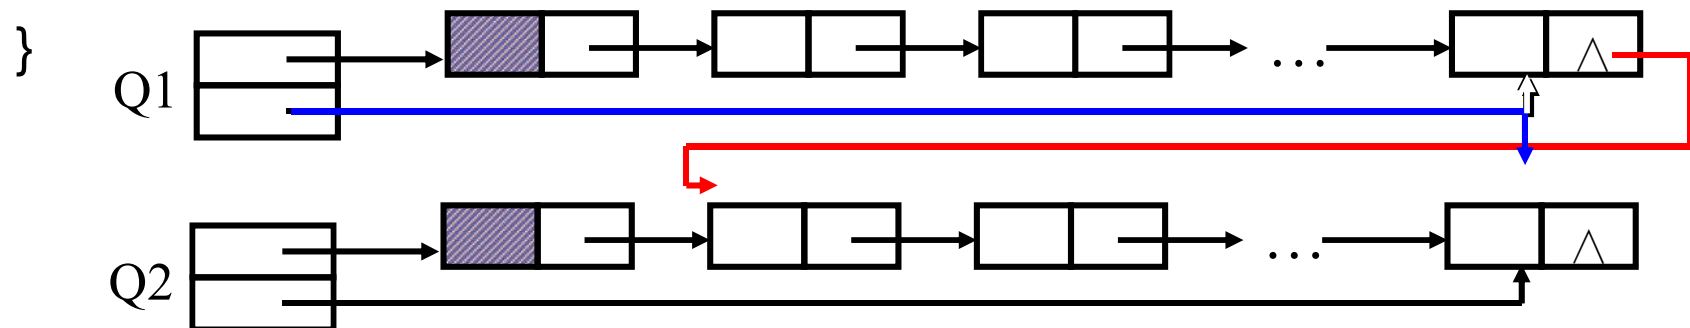
```
void Concatenate(Queue Q1, Queue Q2)
```

```
{ if ( !EMPTY( Q2 ) ) {
```

```
    Q1.rear->next=Q2.front->next;
```

```
    Q1.rear=Q2.rear;
```

```
}
```





Performance of radix sort

Performance analysis

- n : the number of records; d : the number of digits of integer key; r : the radix
- **Time complexity:**
 - Assignment $O(n)$, gathering $O(r)$, d times of assignment and gathering
 - $O(d(n+r))$
- **Space complexity:**
 - The number of head and tail of queue is $2r$
 - The number of link field of each record is n
 - $O((n+r))$



Comparison of Sorting Algorithms

- Performance factors:
 - (1)Running time;
 - (2)Space ;
 - (3)Stability ;
 - (4)Simple ;
 - (5)Input distribution
 - (6) Number of records to be sorted



Comparison of Running Time

Sorting algorithms	Average case	Best case	Worst case
Insertion sort	$O(n^2)$	$O(n)$	$O(n^2)$
Shell sort	$O(n \log_2 n)$	$O(n^{1.3})$	$O(n^2)$
Bubble sort	$O(n^2)$	$O(n)$	$O(n^2)$
Quick sort	$O(n \log_2 n)$	$O(n \log_2 n)$	$O(n^2)$
Selection sort	$O(n^2)$	$O(n^2)$	$O(n^2)$
Heap sort	$O(n \log_2 n)$	$O(n \log_2 n)$	$O(n \log_2 n)$
Merge sort	$O(n \log_2 n)$	$O(n \log_2 n)$	$O(n \log_2 n)$
Radix sort	$O(d(n+r))$	$O(d(n+r))$	$O(d(n+r))$



Space and stability

Sorting algorithms	auxiliary space	stability
Insertion sort	$O(1)$	Yes
Shell sort	$O(1)$	No /3,2,2'(d=2,d=1)
Bubble sort	$O(1)$	Yes
Quick sort	$O(\log_2 n) \sim O(n)$	No /2,2',1
Selection sort	$O(1)$	No /2,2',1
Heap sort	$O(1)$	No/1,2,2'(min heap)
Merge sort	$O(n)$	yes
Radix sort	$O(n+r)$	yes



Comparison of Simpleness & record key distribution

1) Simple Algorithms:

- Insertion sort
- Selection sort
- Bubble sort

2) Not simple Algorithms:

- Shell sort
- Heap sort
- Quick sort

Not affected by the record key distribution

- Selection sort
- Heap sort
- Merge sort

Affected by the record key distribution

- Insertion sort
- Bubble sort
- Quick sort



External Sorting



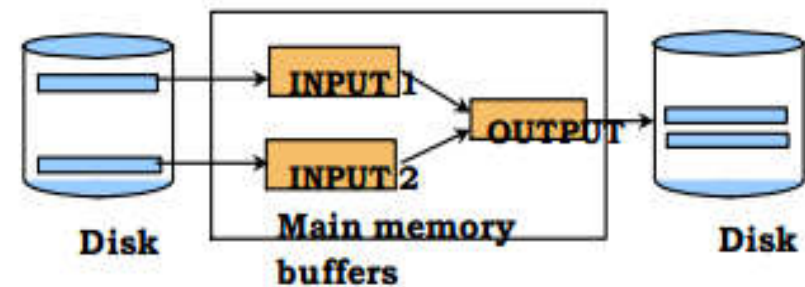
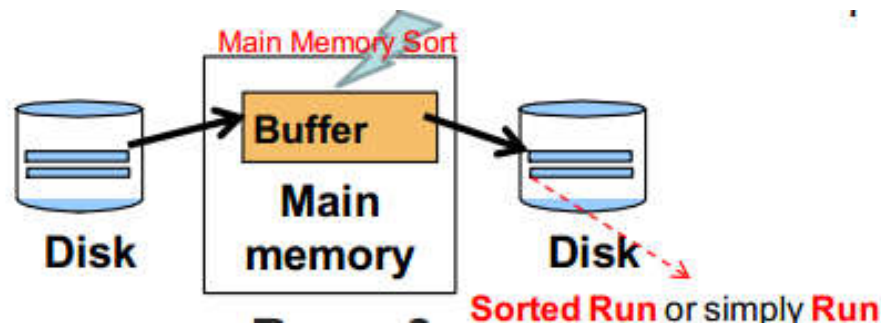
External Sorting Introduction

- **Problem:** We can't sort 1TB of data with 1GB of RAM (i.e., more data than available memory) in main memory
- **Solution:** Utilize an External Sorting Algorithm
 - **External sorting** refers to the sorting of a file that resides on secondary memory (e.g., disk, flash, etc).
 - **Internal sorting** refers to the sorting of an array of data that is in RAM (quick-, merge-, insertion-, selection-, radix-, bucket-, bubble-, heap-, sort algorithms)
- **Objective:** Minimize number of I/O accesses.



Two-Way External Merge-Sort

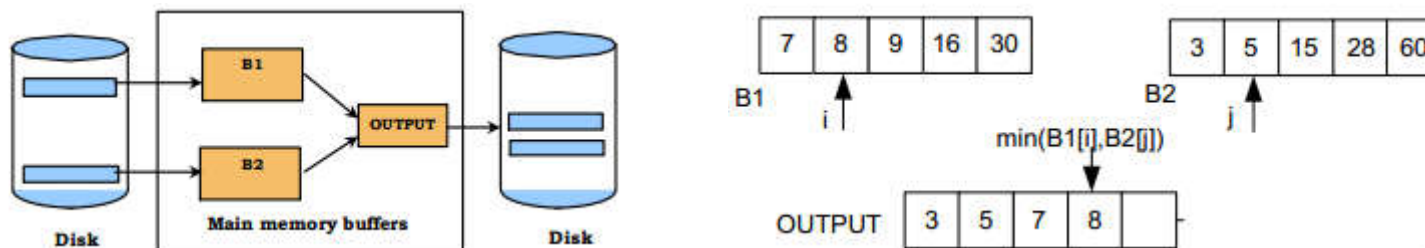
- Let us consider the simplest idea for external sorting
 - Assumption: Only 3 Buffer pages are available
 - Idea: Divide and Conquer (similarly to MergeSort, Quicksort)
- Idea Outline
 - Pass 0 (Sort Lists): For every page, read it, sort it, write it out
 - Only one buffer page is used!
 - Now we need to merge them hierarchically
 - Pass 1, 2, ..., etc. (Merge Lists):
 - For this step we need three buffer pages!
 - see next page for merging concept





Merging Lists

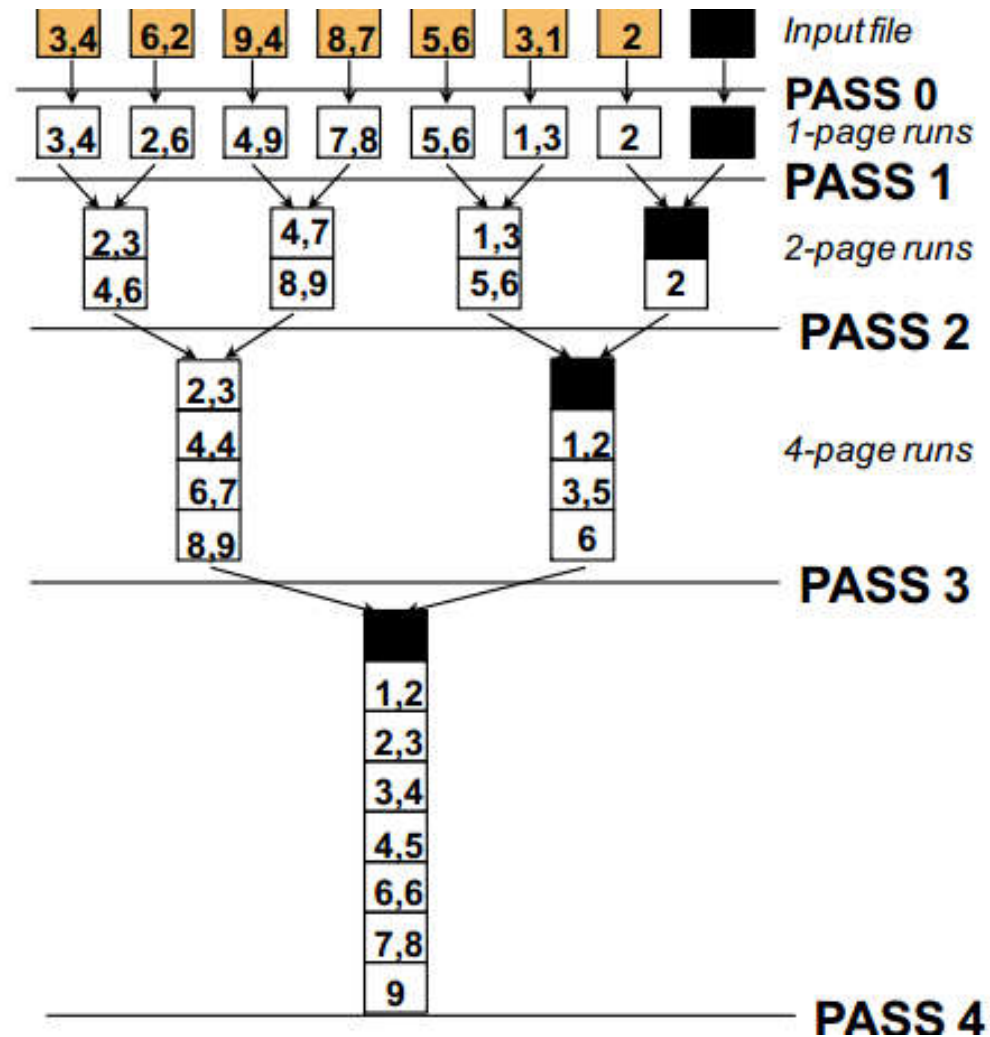
- Merging Lists Outline (Phase 1,2,...)
 1. **Load** the next sorted *runs* R1 and R2 into main memory buffers B1 and B2 a page-at-a-time (i.e., initially first page from each run) (see left figure)
 - Obviously $R1 \geq B1$ and $R2 \geq B2$ (a *Run* might be larger than a Buffer)
 - The rest pages will be loaded to main memory during subsequent steps.
 2. Initialize indices i, j to the head of each list (i.e., $i=j=0$)
 3. **Compare** $B1[i]$ with $B2[j]$ and **move the smallest item to OUTPUT buffer**.
 - If $B1[i]$ was smallest item then $i++$ else $j++$ (see right figure)
 - If OUTPUT gets full, it is appended to the end of a file on DISK and cleared in RAM.
 4. **Repeat** the above until either index i or j reaches the end of its buffer.
 - At this point write the remaining records to OUTPUT, flush it to disk and finish.
 5. Repeat procedure from 1-4 until all *runs* have been traversed.





Two-Way External Merge Sort Example

**Example
Execution**





Cost of Two-Way External Merge Sort

- Each pass we read + write each of N pages in file. **Pass 0**

- Number of passes:
$$= \lceil \log_2 N \rceil + 1$$

e.g., for $N=7$, $N=5$ and $N=4$

$$\lceil \log_2 7 \rceil + 1 = \left\lceil \frac{\log_{10} 7}{\log_{10} 2} \right\rceil + 1 = \lceil 2.8 \rceil + 1 = 4$$

$$\lceil \log_2 5 \rceil + 1 = \lceil 2.3 \rceil + 1 = 4$$

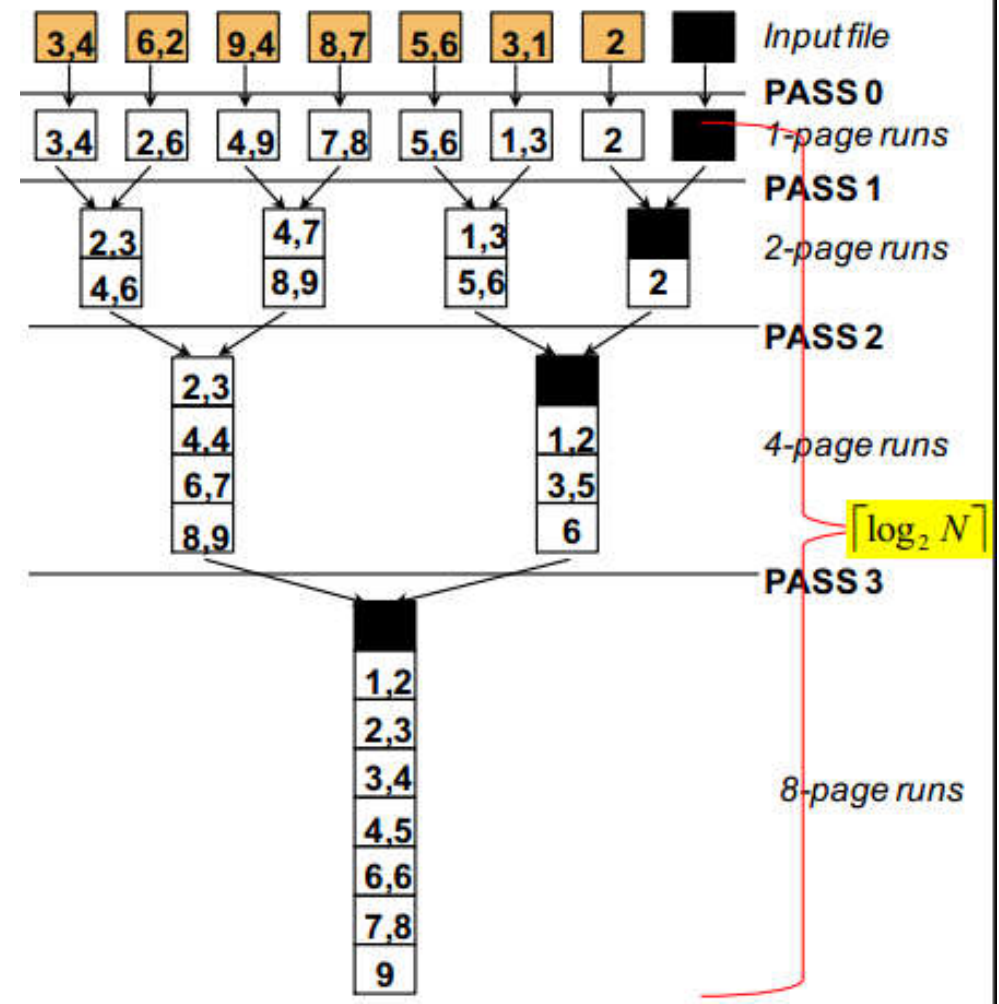
$$\lceil \log_2 4 \rceil + 1 = \lceil 2 \rceil + 1 = 3$$

- Total (I/O) cost is:
$$2N * (\# \text{ passes})$$

e.g., for $N=7$

$$2 * 7 * (\lceil \log_2 7 \rceil + 1) = 2 * 7 * 4 = 56$$

- i.e., (read+write) * 7 pages * 4 passes
- That can be validated on the right figure
 - Pass#0=2*7 Pass#1=2*7
 - Pass#2=2*7 Pass#3=2*7





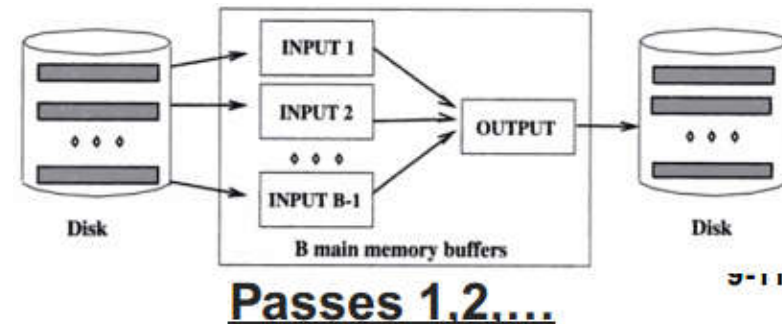
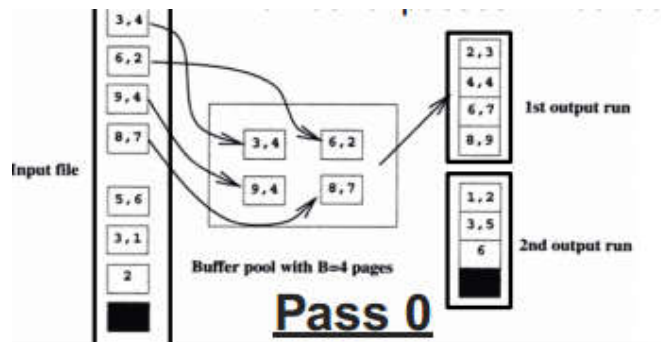
Internal Sorting Remarks

- External Sorting Algorithms utilize some Internal Sorting Algorithm to sort records in main memory.
- What In-Memory Algorithms do real DBMSes use?
 - Oracle uses InsertionSort
 - Microsoft uses MergeSort
 - IBM uses RadixSort (quite similar to BucketSort)



General External Merge Sort

- Let's turn the 2-way Merge-sort Algorithm into a Practical Algorithm.
 - Assumption: B Buffer pages are available
 - Idea: Merge (B-1) pages in each step rather than only 2 (faster!)
- Idea Outline
 - Pass 0 (Sort): Sort the N pages using B buffer pages
 - Use B buffer pages for input
 - That generates $N_1 = \lceil N/B \rceil$ sorted runs e.g., N=8 and B=4 $\Rightarrow N_1=2$
 - Pass 1, 2, ..., etc. (Merge): Perform a (B-1)-way merge of runs
 - Use (B-1) buffer pages for input, 1 page for output
 - Number of passes will be reduced dramatically!





Cost of External Merge Sort

- | | (B-1)-way Sort | 2-way Sort |
|---------------------|--|------------------------------|
| • Number of passes: | $1 + \lceil \log_{B-1} \lceil N/B \rceil \rceil$ | $1 + \lceil \log_2 N \rceil$ |
| • I/O Cost = | $2N * (\text{\# of passes})$ | $2N * (\text{\# of passes})$ |

Example: $N=108$ pages and $B=5$ Buffer pages

- **#Passes** = $1 + \lceil \log_4 \lceil 108/5 \rceil \rceil = 4$
- **IO Cost** = Read + Write 108 pages for 4 passes $\Rightarrow 2 * 108 * 4 = 864$ Ios
 - **Pass 0:** $108/5 = 22$ sorted runs of 5 pages each (last run is only 8 pages)
 - **Pass 2:** $22/4 = 2$ sorted runs, 80 pages ($4 * 20$ pages) and 28 pages
 - **Pass 3:** Sorted file of 108 pages



Number of Passes of External Sort

- External Merge Sort is quite efficient!
- With only $B=257$ (~1MB) Buffer Pages we can sort $N=1$ Billion records with four (4) passes ... in practice B will be larger
- Two-Way Mergesort would require $\lceil \log_2 10^9 \rceil + 1$ passes!

N	B=3	B=5	B=9	B=17	B=129	B=257
100	7	4	3	2	1	1
1,000	10	5	4	3	2	2
10,000	13	7	5	4	2	2
100,000	17	9	6	5	3	3
1,000,000	20	10	7	5	3	3
10,000,000	23	12	8	6	4	3
100,000,000	26	14	9	7	4	4
1,000,000,000	30	15	10	8	5	4

Results generated with formula: $1 + \lceil \log_{B-1} \lceil N/B \rceil \rceil$



Sorting algorithms visualization



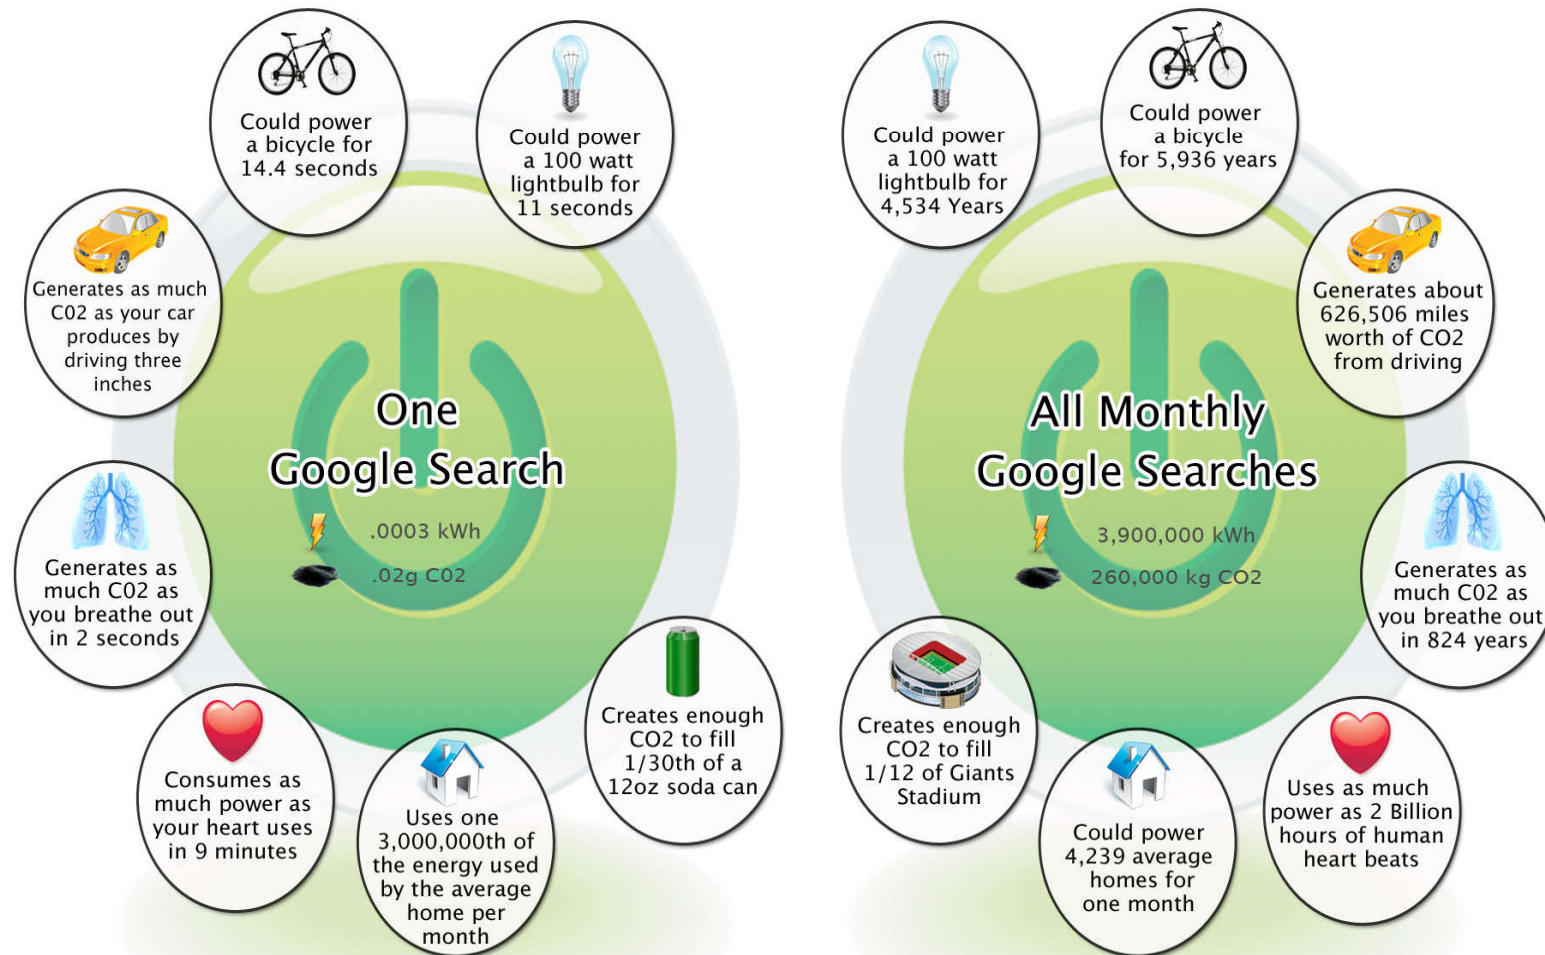
Map/Reduce



The Energy

Used in Google Search

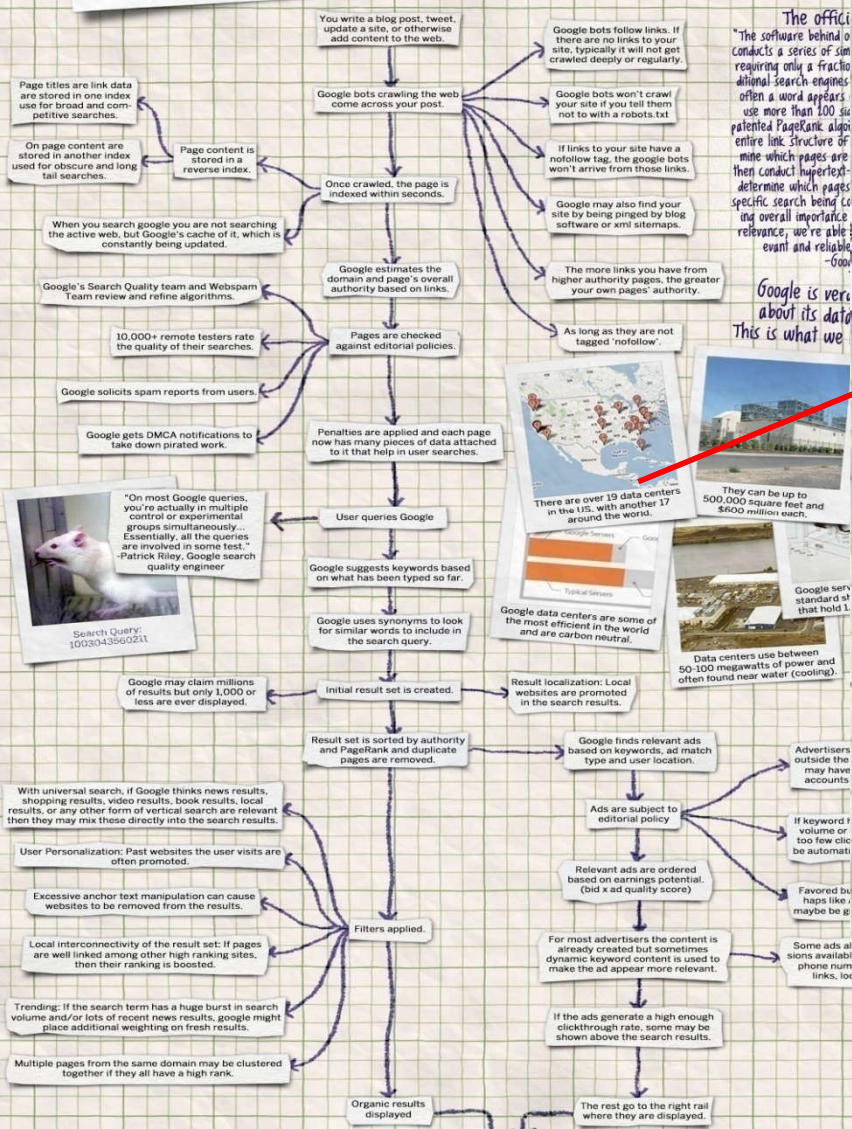
As the world's largest search engine, Google processes nearly 13 Billion monthly searches. They are able to handle such large volumes of data because they have huge datacenters with thousands of servers capable of handling immense capacities. Such large amounts of computing power require a great deal of electricity. This electricity consumption translates directly into carbon emissions. Take a look below:



Google (graphic)

How google works

Google Search I'm Feeling Lucky



and this is all done in less than a second, 300 millions times a day generating over \$20 billion a year for Google!

The official software behind Google conducts a series of simulations requiring only a fractional search engines often a word appears. use more than 100 simulated link structure of entire link structure of mine which pages are then conduct hyper-text-determine which pages specific search being to ing overall importance relevance, we're able event and reliable

Google is very about its data This is what we

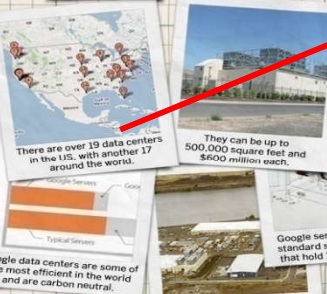
higher authority pages, the greater your own pages' authority.

As long as they are not tagged 'nofollow'.

Google is very secretive about its data centers. This is what we know so far.



There are over 19 data centers in the US, with another 17 around the world.



Google data centers are some of the most efficient in the world and are carbon neutral.



Data centers use between 50-100 megawatts of power and often found near water (cooling).

Google data centers are some of the most efficient in the world and are carbon neutral.

They can be up to 500,000 square feet and \$600 million each.

Google servers are housed in standard shipping containers that hold 1,160 servers each.

Data centers use between 50-100 megawatts of power and often found near water (cooling).

and this is all done in less than a second, 300 millions times a day generating over \$20 billion a year for Google!



Cloud Computing:

Google five core technologies



- GFS (2003)
 - Google File System
- MapReduce (2004)
 - Simplified data processing on large cluster
- BigTable (2006)
 - A distributed storage system for structured data
- Chubby (2006)
 - Lock service for loosely-coupled distributed systems
- BORG
 - Cluster management and schedule



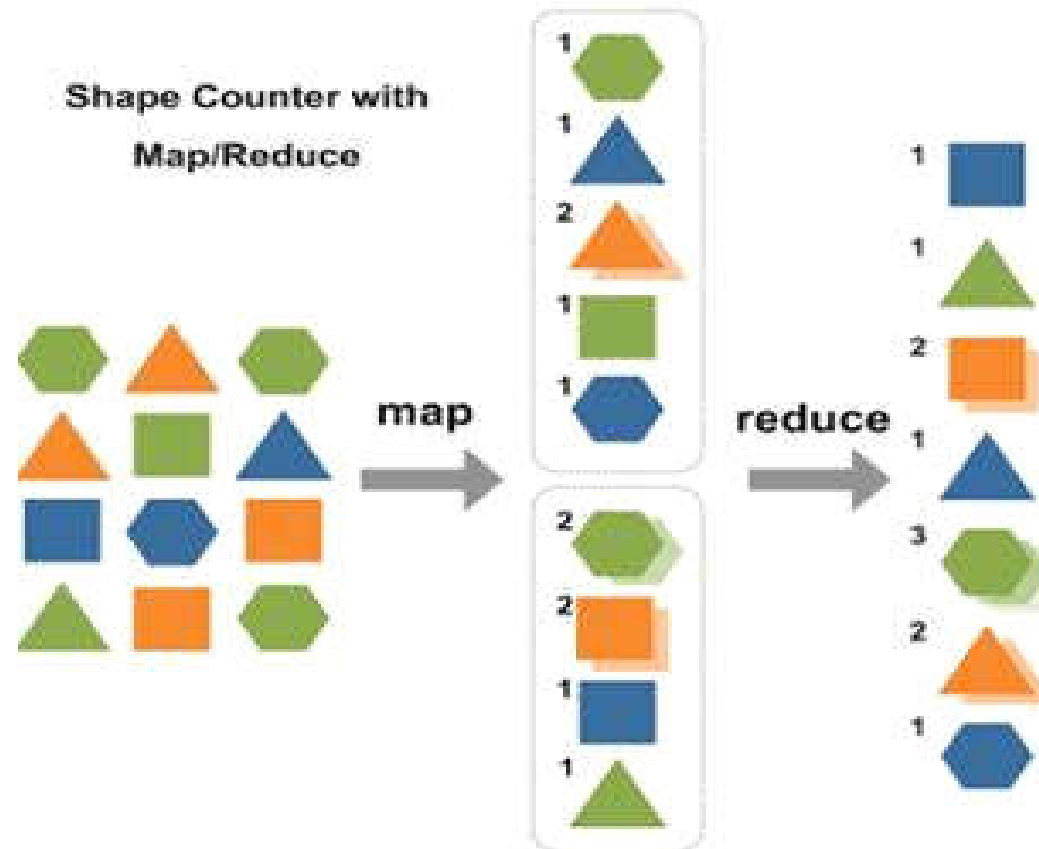
Concurrent & distributed programming model: Map/Reduce

- A programming paradigm
 - Borrows from functional programming
 - Users implement interface of two functions: map()/reduce()
- For processing large data set (> 1 TB)
- Exploits large set of commodity computers (Clusters: +1k CPUs)
- Executes process in distributed manner
- Offers high degree of transparencies
- In other words:
 - Simple, easy, and maybe suitable for your tasks !!!





Simple example: Shape Counter





MapReduce: Map

*Map (input_key, input_value) =>
(output_key, intermediate_value) list*

- Records from the data source are fed into the map function as (key, value) pairs:
 - data source: lines out of files, rows of a database, etc
 - (key, value) of data source: (filename, offset)
- **map()** produces one or more intermediate values along with an output key from the input.
 - {<key1,1>, <key2,1>, <key3,1>, <key4,1>}



MapReduce: Reduce

*Reduce (output_key, intermediate_value list) =>
output list // usually only one final value per key*

- After the map phase is over, all the intermediate values for a given output key are combined together into a list
 - <key → 1, 1, 1>
- *reduce()* combines those intermediate values into one or more final values for that same output key
 - in practice, usually only one final value per key
 - <key, 3>



Two phases of data processing

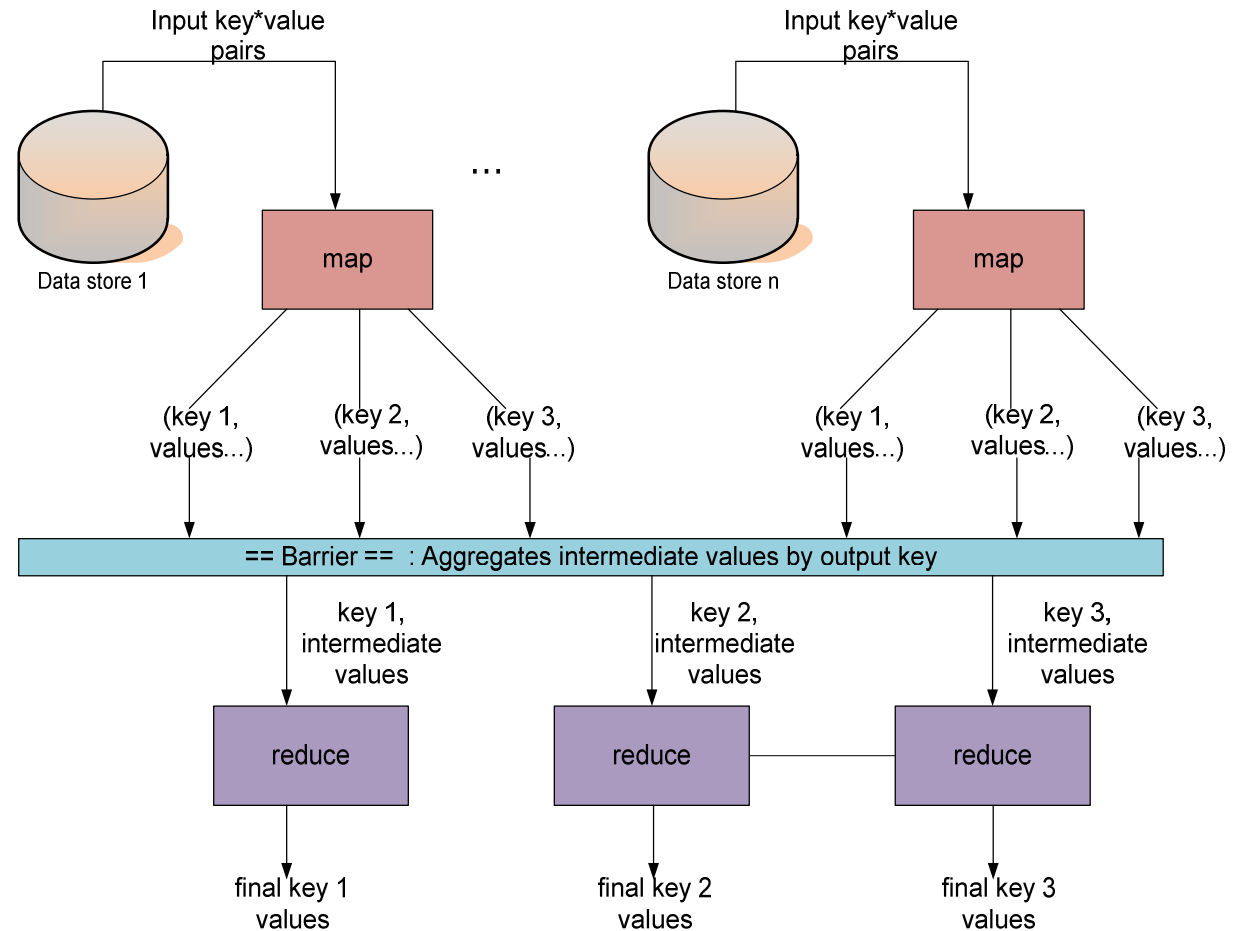
Map (*in_key*, *in_value*) \rightarrow $\{(key_j, value_j) \mid j = 1 \dots k\}$

//identify & combine

Aggregates intermediate values by output key

Reduce (*key*, [*value*₁, ... *value*_{*m*}]) \rightarrow (*key*, *f_value*)

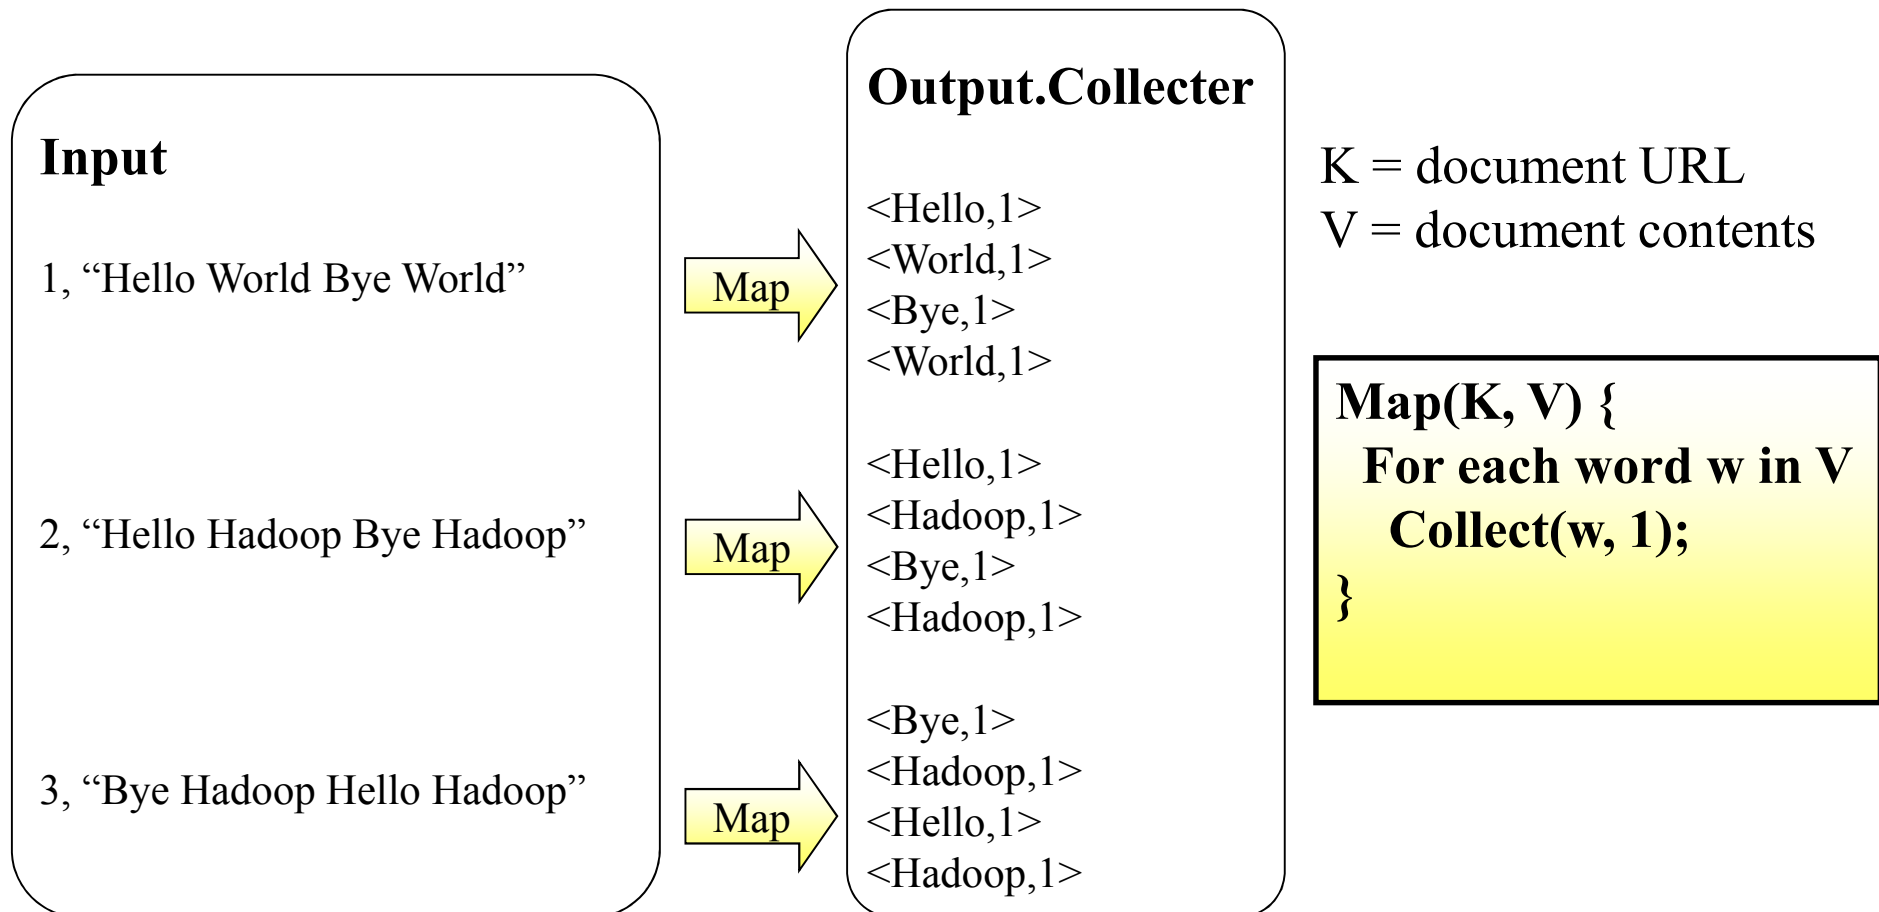
// group





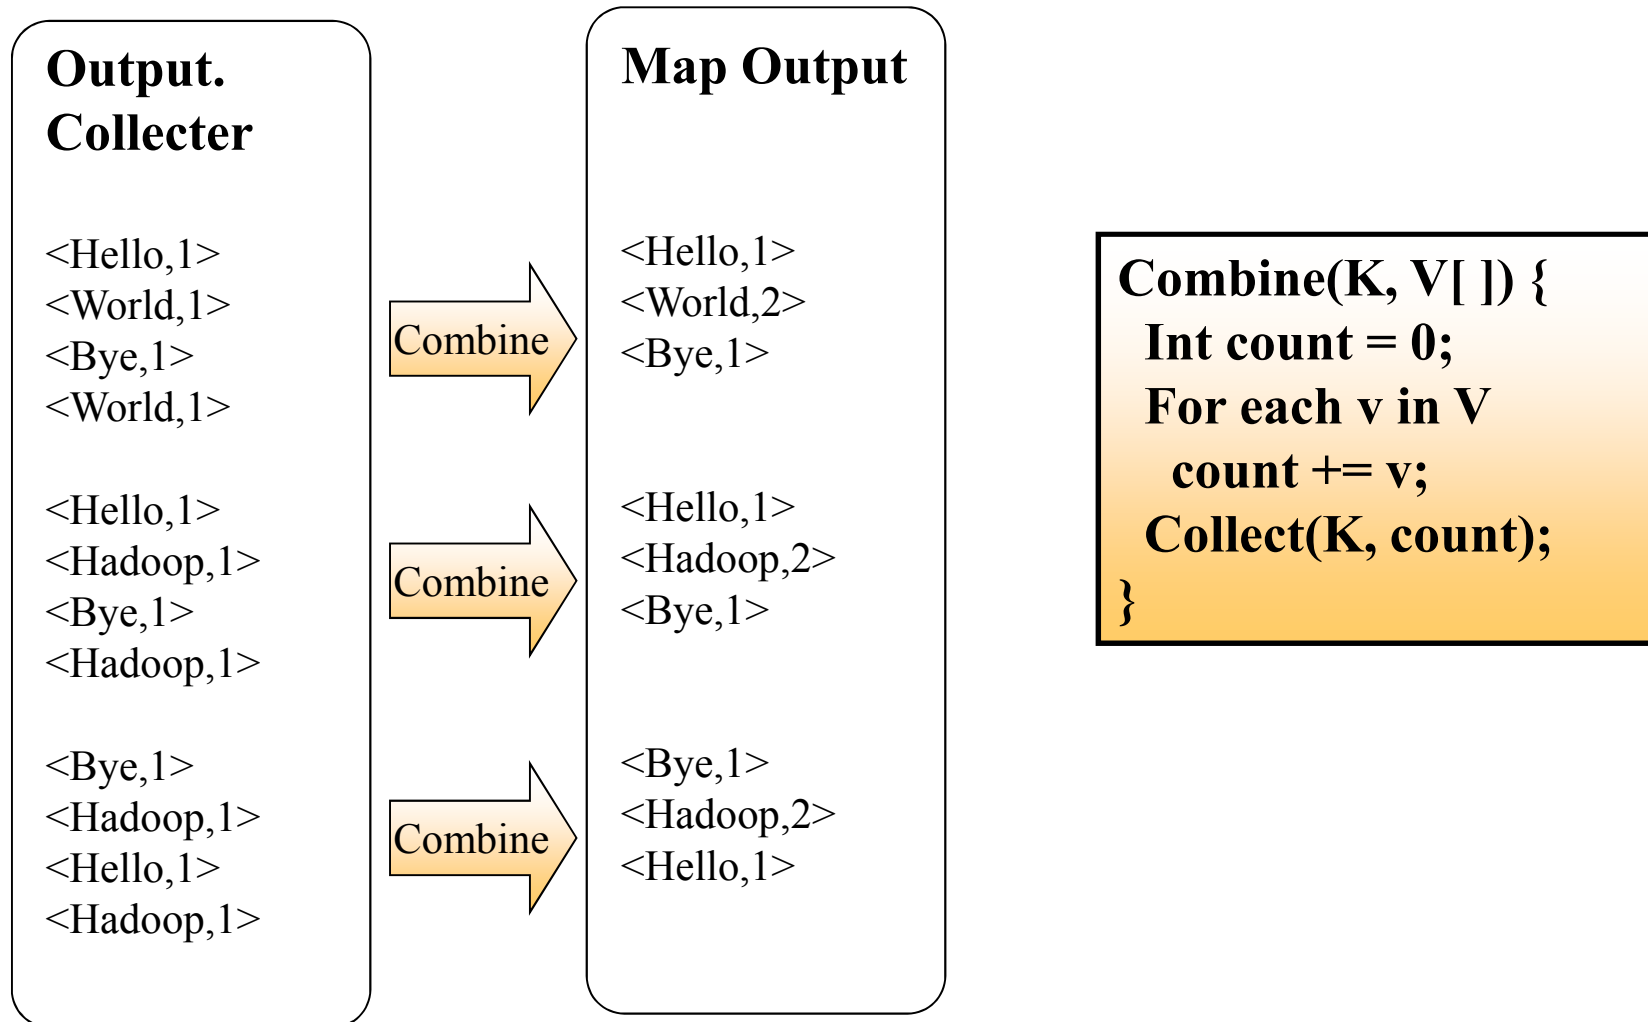
Example: WordCount 1/4

$Map(in_key, in_value) \rightarrow \{(key_j, value_j) \mid j = 1 \dots k\}$





Example: WordCount 2/4





Example: WordCount 3/4

$Reduce(key, [value_1, \dots, value_m]) \rightarrow (key, f_value)$

Reduce Input

<Hello,1>
<World,2>
<Bye,1>

<Hello,1>
<Hadoop,2>
<Bye,1>

<Bye,1>
<Hadoop,2>
<Hello,1>

Internal Grouping

<Bye \rightarrow 1, 1, 1>

<Hadoop \rightarrow 2, 2>

<Hello \rightarrow 1, 1, 1>

<World \rightarrow 2>

Reduce

Reduce

Reduce

Reduce

```
Reduce(K, V[ ]) {  
  Int count = 0;  
  For each v in V  
    count += v;  
  Collect(K, count);  
}
```

Reduce Output

<Bye, 3>
<Hadoop, 4>
<Hello, 3>
<World, 2>



Example: WordCount 4/4

Input:

File containing words

Hello World Bye World
Hello Hadoop Bye Hadoop
Bye Hadoop Hello Hadoop

MapReduce

Output:

Number of occurrences
of each word

Bye 3
Hadoop 4
Hello 3
World 2



Based on Functional Programming

- Functional operations do not modify data structures
 - They always create new ones
- Original data still exists in unmodified form
- Data flows are implicit in program design
- Order of operations does not matter



More Challenges (1/2)

- Automatic parallelization & distribution
 - All values are processed independently
 - Bottleneck: reduce phase can't start until map phase is completely finished.
- Fault-tolerance: nice retry/failure semantics
 - master/slave: periodical ping/checkpoint
 - re-execution
 - file chunks are automatically distributed and replicated
- Locality
 - mappers on same machine as physical file data



More Challenges (2/2)

- Optimization
 - 64MB blocks: same size as GFS chunks
 - mini-reduce phase to occur for saving bandwidth
 - Combiners can run on same machine as a mapper
- Automatic load balancing & high throughput
- I/O scheduling
- Monitoring & status updates



Example applications: 1. Search

- **Input:** (lineNumber, line) records
- **Output:** lines matching a given pattern
- **Map:**

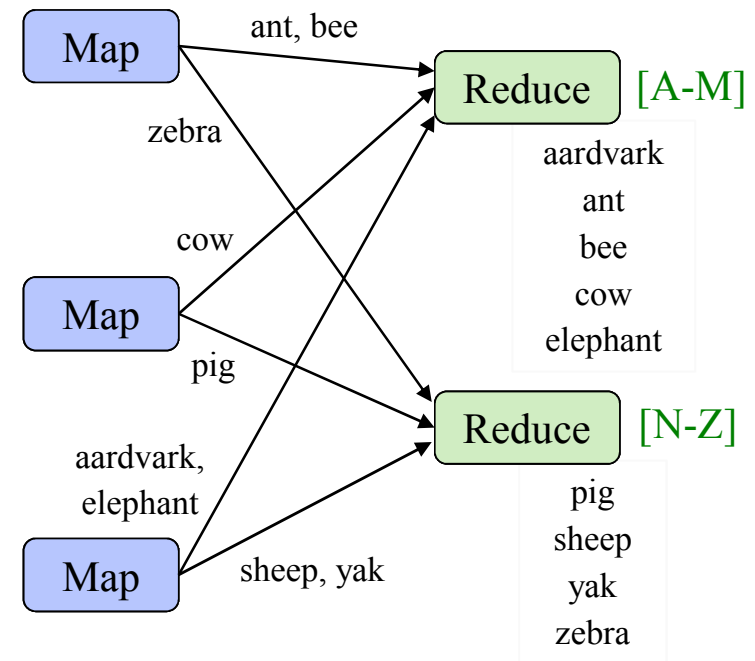
```
if(line matches pattern):  
    output(line)
```
- **Reduce:** identify function
 - Alternative: no reducer (map-only job)



Example applications: 2. Sort

- **Input:** (key, value) records
- **Output:** same records, sorted by key

- **Map:** identity function
- **Reduce:** identify function
- **Trick:** Pick partitioning function h such that $k_1 < k_2 \Rightarrow h(k_1) < h(k_2)$





Example applications: 3. Inverted Index



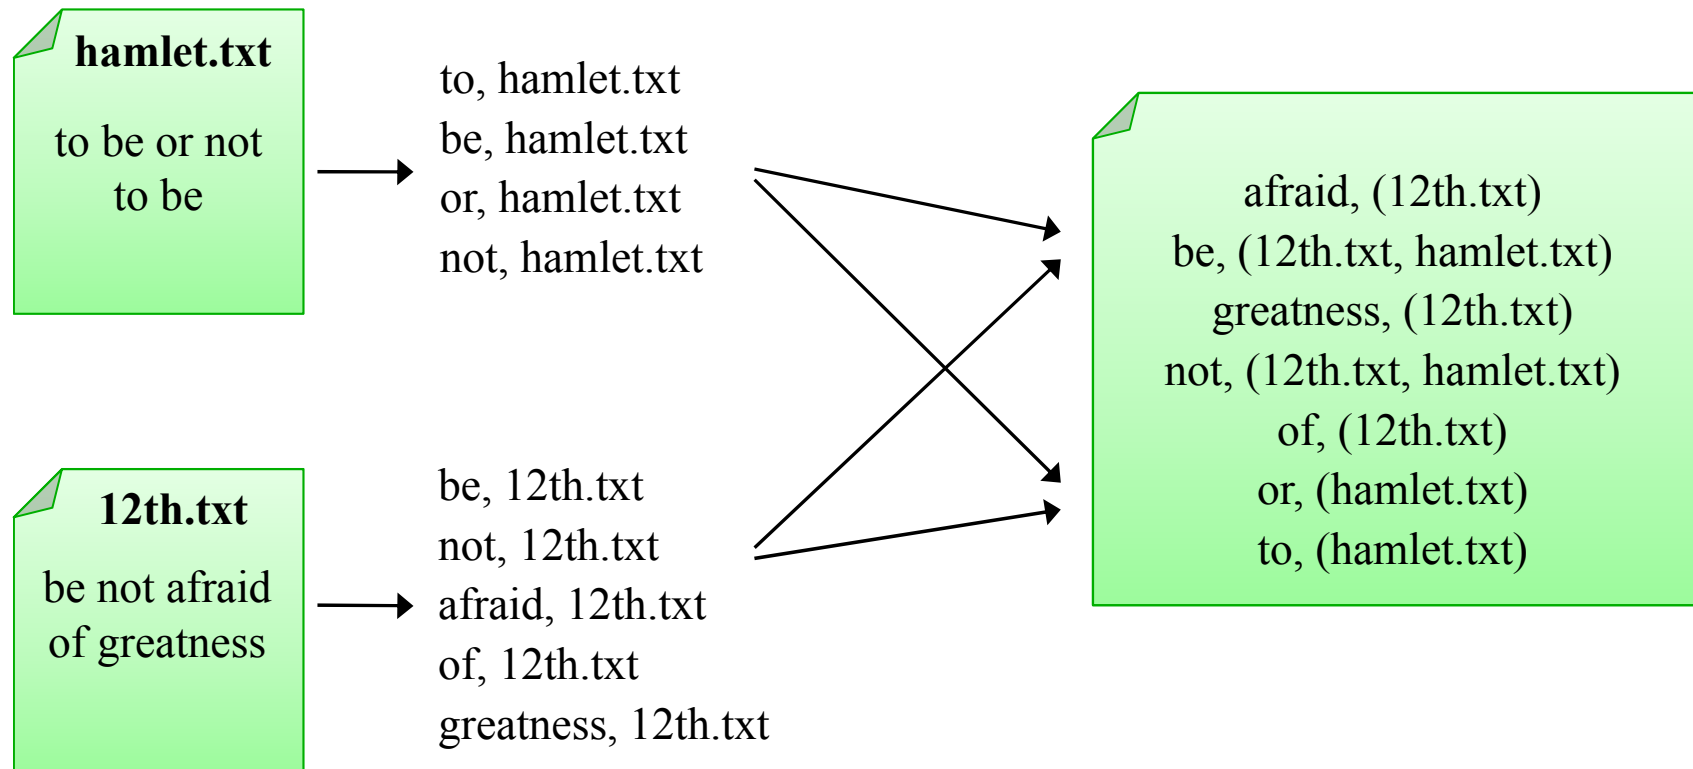
- **Input:** (filename, text) records
- **Output:** list of files containing each word
- **Map:**

```
foreach word in text.split():  
    output(word, filename)
```
- **Combine:** unify filenames for each word
- **Reduce:**

```
def reduce(word, filenames):  
    output(word, sort(filenames))
```



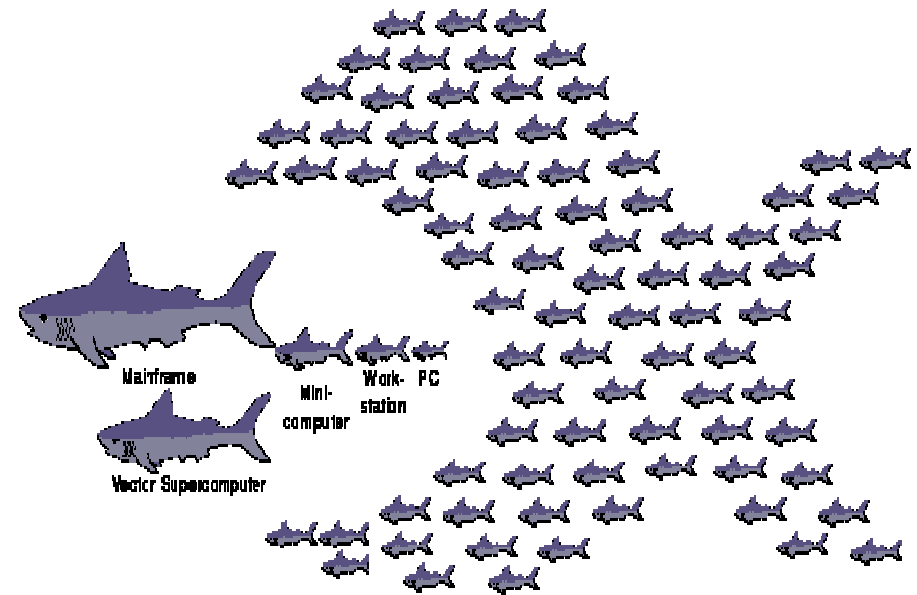
Inverted Index Example





Summary

- A simple programming model for processing large dataset on large set of computer cluster
- Clean abstraction for large scale data processing
- Very robust in fault tolerance and error handling
- Fun to use,
 - focus on problem, and
 - let the library deal with the messy detail





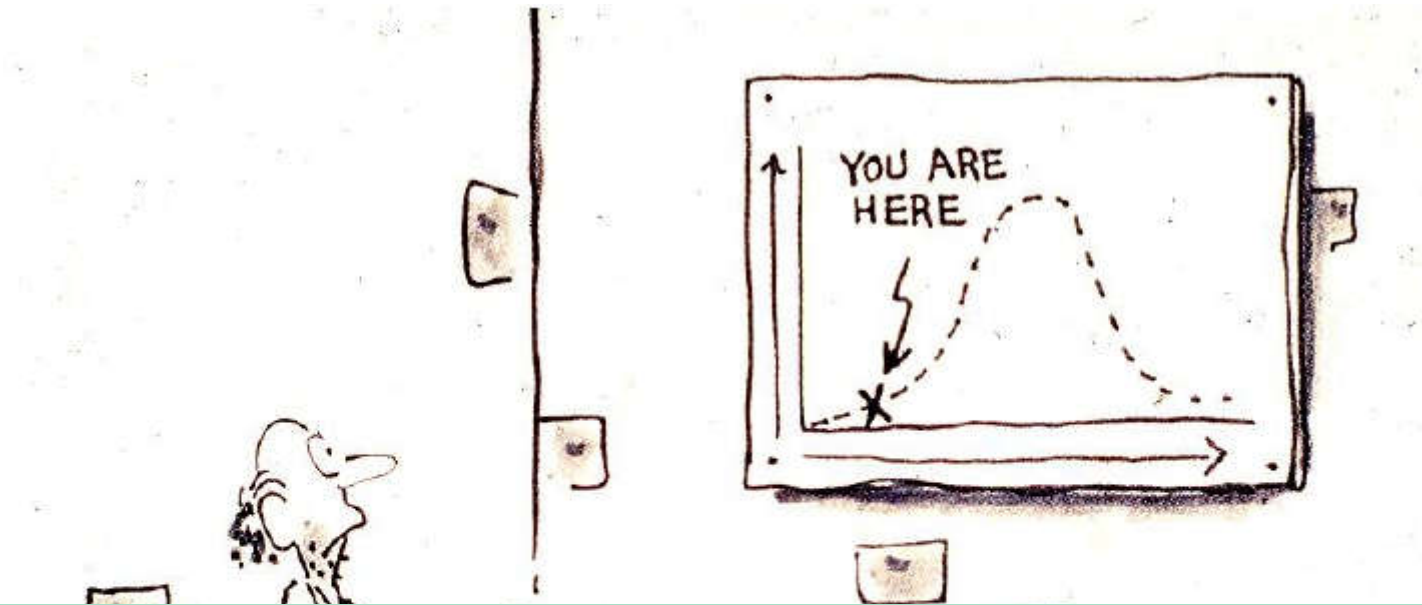
"Successful people are not gifted; they just work hard, then succeed on purpose." —
G. K. Nielson

"Scientists investigate that which already is; Engineers create that which has never
been." — Albert Einstein.

" **Computers do not solve problems, they execute solutions** " — Laurent Gasser

"The real problem is not whether machines think but whether men do." — B. F.
Skinner

"Your beliefs become your thoughts, Your thoughts become your words, Your
words become your action, Your actions become your habits, Your habits
become your values, You values become your destiny." — Mahatma Gandhi



*You make the world different.
Never give up, Never!*